

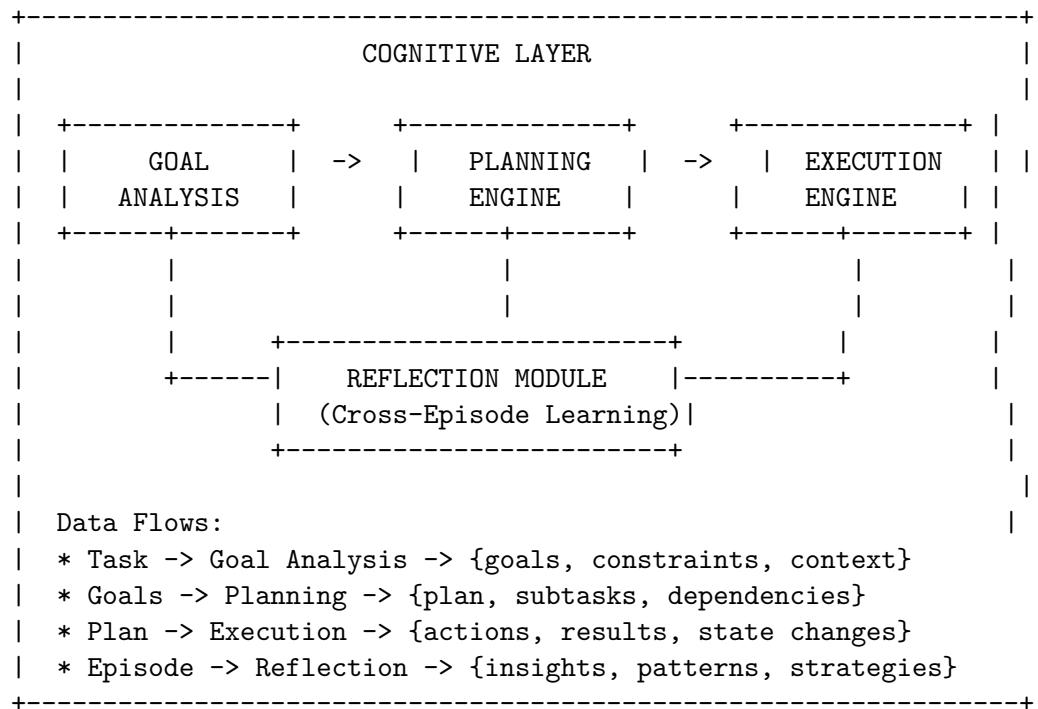
# Contents

<b>Cognitive Layer: Core Reasoning Components</b>	<b>2</b>
Overview . . . . .	2
1. Goal Analysis Module . . . . .	2
1.1 Purpose & Responsibilities . . . . .	2
1.2 Architecture Interna . . . . .	2
1.3 Semantic Parser . . . . .	4
1.4 Goal Extractor . . . . .	5
1.5 Constraint Analyzer . . . . .	7
1.6 Context Assessor . . . . .	9
1.7 Complexity Classifier . . . . .	10
1.8 Goal Analysis Output Schema . . . . .	12
2. Planning Engine . . . . .	13
2.1 Purpose & Responsibilities . . . . .	13
2.2 Architecture Interna . . . . .	14
2.3 Strategy Selector . . . . .	15
2.4 Task Decomposer (HTN Focus) . . . . .	16
2.5 Dependency Analyzer . . . . .	18
2.6 Resource Estimator . . . . .	21
2.7 Plan Optimizer . . . . .	23
2.8 Contingency Planner . . . . .	25
2.9 Planning Engine Output Schema . . . . .	28
3. Execution Engine . . . . .	30
3.1 Purpose & Responsibilities . . . . .	30
3.2 Architecture Interna . . . . .	30
3.3 Execution Coordinator . . . . .	31
3.4 Task Executor . . . . .	33
3.5 State Manager . . . . .	35
3.6 Monitoring System . . . . .	37
3.7 Verification Layer . . . . .	39
3.8 Adaptation Engine . . . . .	42
3.9 Execution Engine Output Schema . . . . .	46
4. Reflection Module . . . . .	48
4.1 Purpose & Responsibilities . . . . .	48
4.2 Architecture Interna . . . . .	48
4.3 Episode Analyzer . . . . .	49
4.4 Pattern Extractor . . . . .	51
4.5 Strategy Evaluator . . . . .	54
4.6 Performance Analyzer . . . . .	56
4.7 Knowledge Distiller . . . . .	58
4.8 Memory Updater . . . . .	60
4.9 Reflection Output Schema . . . . .	62
5. Cognitive Layer Integration . . . . .	64
5.1 Complete Flow Example . . . . .	64
5.2 Caratteristiche del Cognitive Layer . . . . .	67
5.3 Caratteristiche di Performance . . . . .	67

# Cognitive Layer: Core Reasoning Components

## Overview

Il Cognitive Layer è il cuore dell'agente - dove avviene il reasoning di alto livello. Comprende 4 moduli interconnessi che implementano il ciclo completo di problem-solving: comprensione del problema, pianificazione, esecuzione, e apprendimento dall'esperienza.



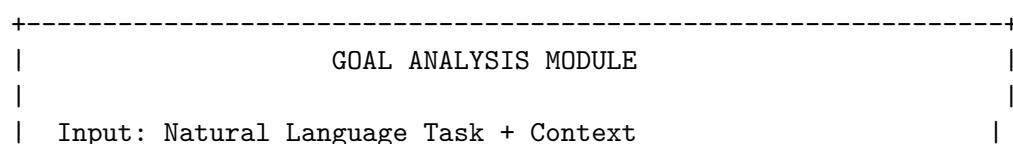
## 1. Goal Analysis Module

### 1.1 Purpose & Responsibilities

**Core Function:** Trasformare task in input (spesso ambiguo, natural language) in rappresentazione strutturata e computazionale.

**Responsibilities:** 1. **Parsing:** Estrarre semantica da input naturale 2. **Goal Extraction:** Identificare obiettivi primari e secondari 3. **Constraint Identification:** Rilevare vincoli impliciti ed esplicativi 4. **Context Assessment:** Valutare informazioni contestuali rilevanti 5. **Complexity Classification:** Determinare classe problema e strategia appropriata

### 1.2 Architecture Interna



```

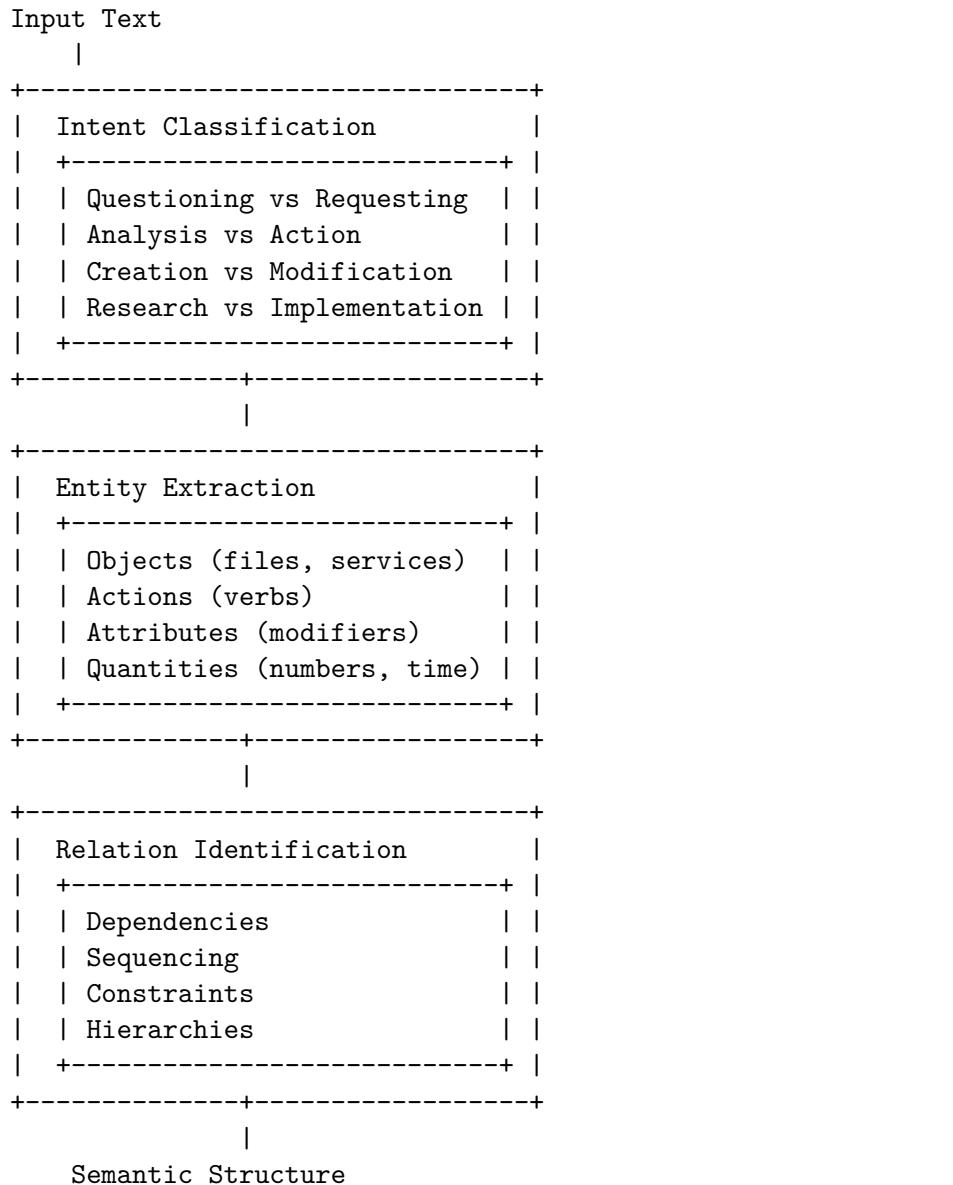
|   |
| +-----+
| | 1. SEMANTIC PARSER           |
| | * Intent classification    |
| | * Entity extraction         |
| | * Relation identification   |
| +-----+
|   |
| +-----+
| | 2. GOAL EXTRACTOR          |
| | * Primary goal identification|
| | * Sub-goal decomposition    |
| | * Success criteria extraction|
| +-----+
|   |
| +-----+
| | 3. CONSTRAINT ANALYZER    |
| | * Explicit constraints     |
| | * Implicit constraints (from context)|
| | * Safety bounds            |
| | * Resource limits          |
| +-----+
|   |
| +-----+
| | 4. CONTEXT ASSESSOR        |
| | * Retrieve relevant history|
| | * Identify domain/environment|
| | * Assess available resources|
| +-----+
|   |
| +-----+
| | 5. COMPLEXITY CLASSIFIER   |
| | * Estimate problem difficulty|
| | * Classify into taxonomy category|
| | * Select appropriate strategy|
| +-----+
|   |
| Output: Structured Goal Representation
{
  primary_goal: {...},
  sub_goals: [...],
  constraints: {...},
  context: {...},
  complexity: "...",
  strategy_hint: "..."
}
+-----+

```

### 1.3 Semantic Parser

**Approach:** Hybrid (LLM reasoning + structured extraction)

**Architecture:**



**Example Flow:**

Input: "Refactor the authentication module to use JWT tokens, ensure backward compatibility, and update tests"

Semantic Parser Output:

```
{  
  intent: "MODIFICATION",  
  entities: {  
    objects: ["authentication module", "JWT tokens", "tests"],  
  }  
}
```

```

        actions: ["refactor", "ensure", "update"],
        constraints: ["backward compatibility"]
    },
    relations: {
        primary_action: "refactor authentication module",
        method: "use JWT tokens",
        constraint: "maintain backward compatibility",
        follow_up: "update tests"
    }
}

```

## 1.4 Goal Extractor

**Purpose:** Convertire semantic structure in goal hierarchy esplicita.

### Goal Representation:

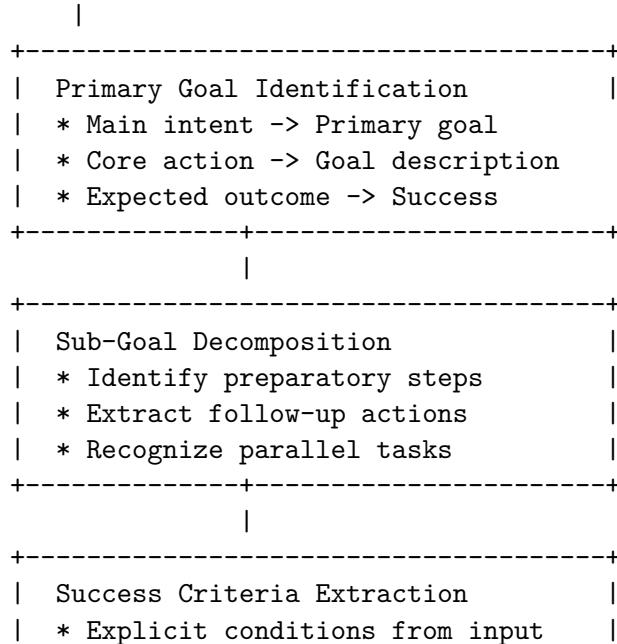
```

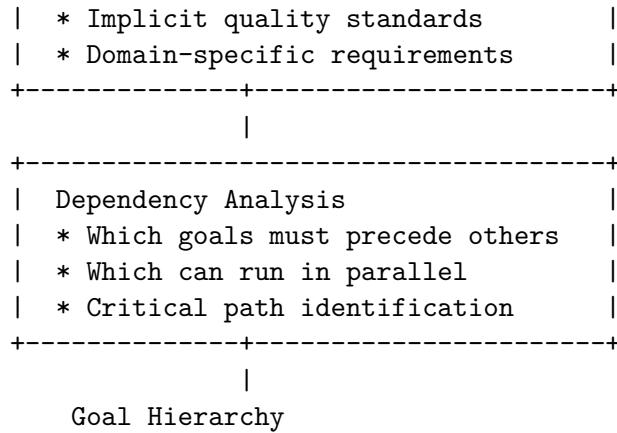
Goal {
    id: unique_identifier,
    type: PRIMARY | SECONDARY | CONSTRAINT,
    description: human_readable,
    success_criteria: [measurable_conditions],
    priority: 1-10,
    dependencies: [goal_ids],
    estimated_complexity: LOW | MEDIUM | HIGH,
    verification_method: how_to_verify_success
}

```

### Extraction Process:

#### Semantic Structure





### **Example Flow:**

Input (from parser):

```
{
  intent: "MODIFICATION",
  entities: {
    objects: ["authentication module", "JWT tokens", "tests"],
    actions: ["refactor", "ensure", "update"],
    constraints: ["backward compatibility"]
  }
}
```

Goal Extractor Output:

```
{
  primary_goal: {
    id: "G1",
    type: "PRIMARY",
    description: "Refactor authentication to use JWT",
    success_criteria: [
      "Authentication uses JWT tokens",
      "All auth flows functional",
      "Tests passing"
    ],
    priority: 10,
    dependencies: [],
    estimated_complexity: "HIGH"
  },
  sub_goals: [
    {
      id: "G1.1",
      type: "SECONDARY",
      description: "Implement JWT token generation",
      success_criteria: ["JWT library integrated", "Token generation working"],
      priority: 9,
      dependencies: [],
      estimated_complexity: "MEDIUM"
    }
  ]
}
```

```

},
{
  id: "G1.2",
  type: "SECONDARY",
  description: "Update authentication flow",
  success_criteria: ["Auth flow uses JWT", "Old flow still works"],
  priority: 8,
  dependencies: ["G1.1"],
  estimated_complexity: "HIGH"
},
{
  id: "G1.3",
  type: "SECONDARY",
  description: "Update test suite",
  success_criteria: ["All tests pass", "New tests for JWT"],
  priority: 7,
  dependencies: ["G1.2"],
  estimated_complexity: "MEDIUM"
}
],
constraints: [
  {
    id: "C1",
    type: "CONSTRAINT",
    description: "Maintain backward compatibility",
    verification_method: "Old clients still authenticate",
    priority: 10
  }
]
}

```

## 1.5 Constraint Analyzer

**Purpose:** Identificare tutti i vincoli che limitano spazio di soluzioni accettabili.

### Constraint Types:

CONSTRAINT TAXONOMY	
1. FUNCTIONAL CONSTRAINTS	
* Must-have features	
* Compatibility requirements	
* Interface contracts	
2. NON-FUNCTIONAL CONSTRAINTS	
* Performance (latency, throughput)	
* Scalability requirements	
* Security requirements	

```

| 3. RESOURCE CONSTRAINTS
|   * Time budget
|   * Cost budget
|   * Token/API call limits
|
| 4. SAFETY CONSTRAINTS
|   * Actions prohibited
|   * Data protection requirements
|   * Approval gates needed
|
| 5. CONTEXTUAL CONSTRAINTS
|   * Environment limitations
|   * Available tools/capabilities
|   * Dependencies on external systems
+-----+

```

### **Analysis Flow:**

```

Input: Task + Context + History
|
+-----+
| Explicit Constraint Extraction      |
| * Parse stated requirements         |
| * Identify "must", "never", etc.    |
+-----+
|
+-----+
| Implicit Constraint Inference     |
| * Domain conventions               |
| * Safety standards                 |
| * Quality expectations             |
+-----+
|
+-----+
| Constraint Validation            |
| * Check for conflicts            |
| * Assess feasibility              |
| * Prioritize if conflicting      |
+-----+
|
+-----+
| Bound Generation                  |
| * Convert to enforceable bounds  |
| * Define verification methods    |
+-----+
|
+-----+
| Constraint Set

```

## 1.6 Context Assessor

**Purpose:** Raccogliere e valutare informazioni contestuali rilevanti per il task.

### Context Dimensions:

CONTEXT ASSESSMENT	
HISTORICAL CONTEXT	
* Similar past tasks	
* Previous failures/successes	
* Learned patterns	
* User preferences	
ENVIRONMENTAL CONTEXT	
* Current working directory	
* Available files/resources	
* System state	
* External service status	
CAPABILITY CONTEXT	
* Available tools	
* API access	
* Model capabilities	
* Resource budgets	
DOMAIN CONTEXT	
* Technical domain (web, ML, data)	
* Language/framework	
* Architecture patterns	
* Best practices	

### Assessment Process:

Task Input	
Query Historical Memory	
* Semantic search similar tasks	
* Retrieve outcomes	

```

| * Extract relevant patterns |
+-----+
|
+-----+
| Assess Environment |
| * Probe file system |
| * Check service availability |
| * Verify resource access |
+-----+
|
+-----+
| Inventory Capabilities |
| * Available tools registry |
| * Model routing options |
| * Budget remaining |
+-----+
|
+-----+
| Identify Domain |
| * Classify technical area |
| * Load domain knowledge |
| * Apply conventions |
+-----+
|
+-----+
Context Structure

```

## 1.7 Complexity Classifier

**Purpose:** Stimare difficoltà problema e selezionare strategia appropriata.

### Classification Dimensions:

```

+-----+
|           COMPLEXITY ASSESSMENT          |
|
| Dimension 1: DECOMPOSITION DEPTH        |
| +- Simple (1 step)                     |
| +- Moderate (2-5 steps)                |
| +- Complex (6-15 steps)                |
| +- Very Complex (>15 steps)           |
|
| Dimension 2: UNCERTAINTY               |
| +- Deterministic (clear path)          |
| +- Low uncertainty (few unknowns)      |
| +- Moderate uncertainty (some exploration needed) |
| +- High uncertainty (research required) |
|
| Dimension 3: INTERDEPENDENCIES         |
| +- Independent (no dependencies)       |
+-----+

```

```

| +- Sequential (linear chain)
| +- Tree (branching dependencies)
| +- Graph (complex interdependencies)
|
| Dimension 4: RISK
| +- No-risk (easily reversible)
| +- Low-risk (can rollback)
| +- Medium-risk (some irreversible)
| +- High-risk (critical operations)
|
| Dimension 5: NOVELTY
| +- Routine (done many times)
| +- Familiar (similar to past tasks)
| +- Novel (first time, but similar domain)
| +- Unprecedented (no prior examples)
+-----+

```

### **Classification Decision Tree:**

```

Task
|
Is it single-step?
+- YES -> SIMPLE
|      +- Routing: Direct execution
|      +- Strategy: Single LLM call
|
+- NO -> Multi-step needed
|
|      Uncertainty level?
|      +- LOW -> STRUCTURED_COMPLEX
|          +- Routing: Planning Engine
|          +- Strategy: HTN planning
|
|      +- HIGH -> EXPLORATORY_COMPLEX
|
|          Risk level?
|          +- LOW/MEDIUM -> ADAPTIVE
|              +- Routing: Planning + Reflection
|              +- Strategy: Iterative refinement
|
|          +- HIGH -> VERIFIED
|              +- Routing: Planning + Verification
|              +- Strategy: Formal verification

```

### **Output Mapping:**

Complexity Classification -> Strategy Selection

SIMPLE:

- No planning phase needed

- Direct execution with single tool call
- Minimal memory usage
- Fast path

**STRUCTURED\_COMPLEX:**

- Hierarchical task decomposition
- Clear dependency management
- Systematic execution
- Standard planning approach

**EXPLORATORY\_COMPLEX:**

- Iterative planning
- Frequent replanning based on discoveries
- Heavy memory usage (episodic)
- Reflection after each major step

**VERIFIED\_COMPLEX:**

- Extensive pre-execution planning
- Verification gates before critical actions
- Human approval for high-risk operations
- Comprehensive logging and traceability

## 1.8 Goal Analysis Output Schema

### Complete Output Structure:

```
GoalAnalysisResult {
    // Core Task Understanding
    task_id: string,
    original_input: string,
    timestamp: datetime,

    // Semantic Understanding
    semantic: {
        intent: Intent,
        entities: Entity[],
        relations: Relation[]
    },

    // Goal Structure
    goals: {
        primary: Goal,
        secondary: Goal[],
        constraints: Constraint[]
    },

    // Context
    context: {

```

```

historical: {
    similar_tasks: Episode[],
    relevant_patterns: Pattern[],
    user_preferences: Preference[]
},
environmental: {
    working_directory: string,
    available_resources: Resource[],
    system_state: State
},
capabilities: {
    available_tools: Tool[],
    model_options: Model[],
    budget: Budget
},
domain: {
    identified_domain: Domain,
    conventions: Convention[],
    best_practices: Practice[]
}
},
// Complexity Assessment
complexity: {
    overall: ComplexityLevel,
    dimensions: {
        decomposition_depth: int,
        uncertainty: UncertaintyLevel,
        interdependencies: DependencyStructure,
        risk: RiskLevel,
        novelty: NoveltyLevel
    },
    recommended_strategy: Strategy
},
// Metadata
confidence: float,
ambiguities: Ambiguity[],
assumptions: Assumption[]
}

```

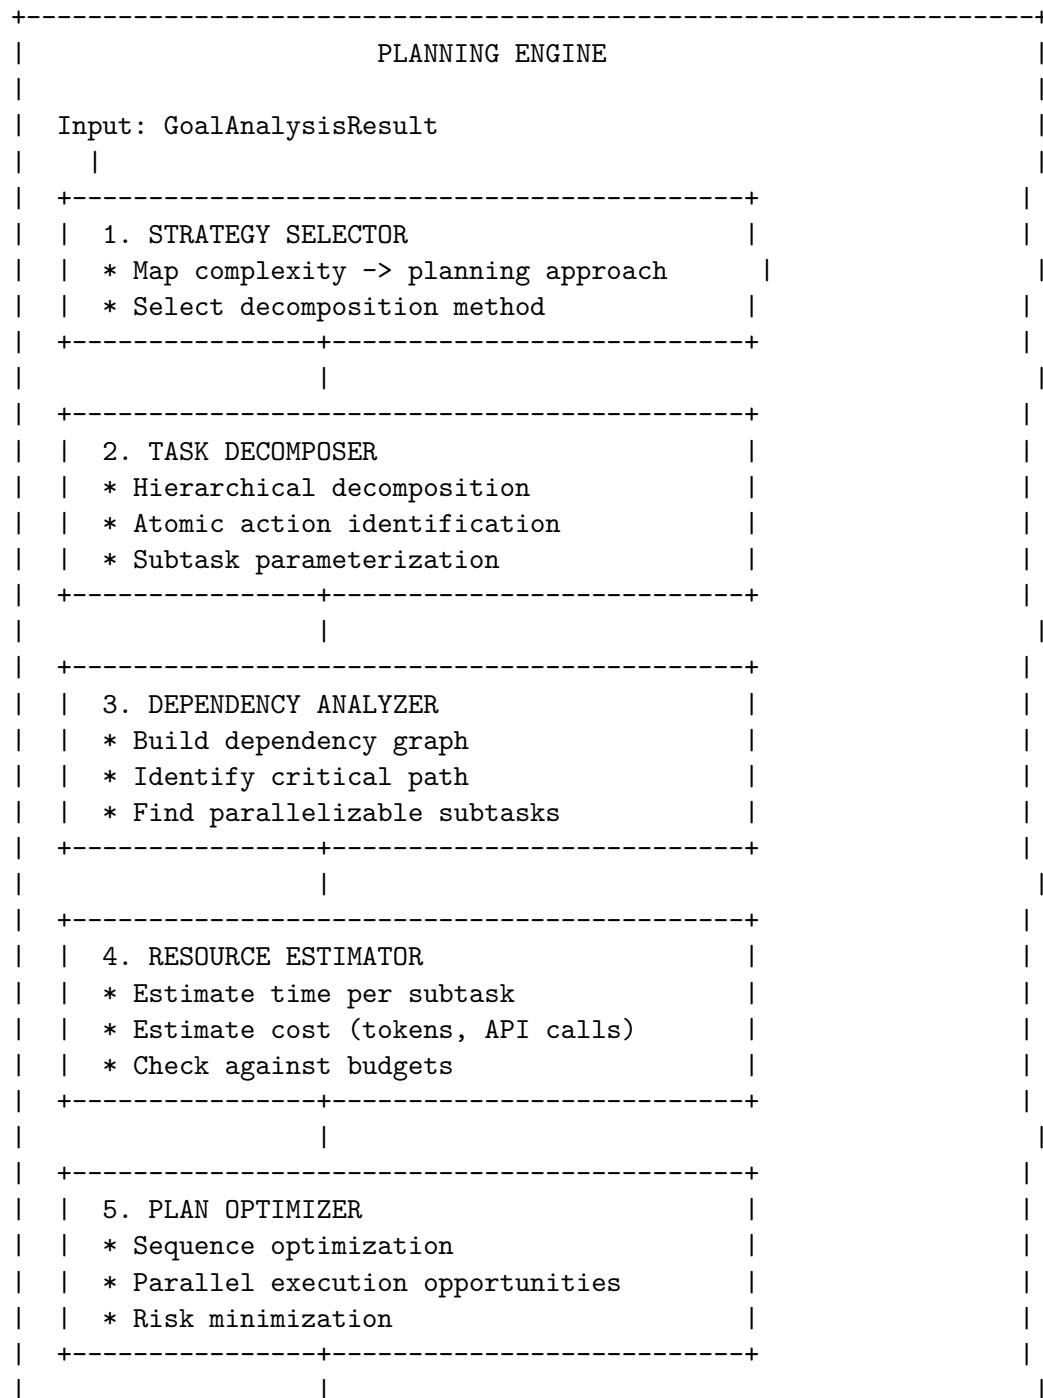
## 2. Planning Engine

### 2.1 Purpose & Responsibilities

**Core Function:** Generare execution plan ottimale che trasforma goal structure in sequenza di azioni concrete ed eseguibili.

**Responsibilities:** 1. **Task Decomposition:** Scomporre goal complessi in subtask atomici 2. **Dependency Management:** Identificare e gestire dipendenze tra subtask 3. **Resource Allocation:** Stimare e allocare risorse (tempo, token, tools) 4. **Strategy Selection:** Scegliere approccio di planning appropriato 5. **Plan Optimization:** Ottimizzare per latenza, costo, o altri obiettivi 6. **Contingency Planning:** Preparare fallback per potenziali failure

## 2.2 Architecture Interna



```

| +-----+
| | 6. CONTINGENCY PLANNER
| | * Identify failure points
| | * Generate fallback strategies
| | * Define recovery procedures
| +-----+
|
| |
| Output: ExecutionPlan
| {
|   subtasks: [...],
|   dependencies: DAG,
|   execution_order: [...],
|   resource_estimates: {...},
|   contingencies: [...]
| }
+-----+

```

## 2.3 Strategy Selector

**Purpose:** Selezionare metodo di planning appropriato basato su caratteristiche task.

### Available Strategies:

```

+-----+
|          PLANNING STRATEGIES
|
| 1. DIRECT EXECUTION
|   When: Simple, single-step tasks
|   Approach: No decomposition, direct tool invocation
|   Overhead: Minimal
|
| 2. LINEAR DECOMPOSITION
|   When: Sequential tasks, clear path
|   Approach: Simple step-by-step breakdown
|   Overhead: Low
|
| 3. HIERARCHICAL TASK NETWORK (HTN)
|   When: Complex tasks, known domain
|   Approach: Recursive decomposition with methods
|   Overhead: Medium
|
| 4. MEANS-ENDS ANALYSIS
|   When: Goal-oriented, reduce difference iteratively
|   Approach: Identify gaps, find operators to reduce
|   Overhead: Medium
|
| 5. CASE-BASED PLANNING
|   When: Similar past tasks available
|   Approach: Retrieve, adapt previous plans
+-----+

```

```

|   Overhead: Low (if cache hit)
|
| 6. EXPLORATORY PLANNING
|   When: High uncertainty, research needed
|   Approach: Iterative: plan -> execute -> learn -> replan
|   Overhead: High
|
| 7. VERIFIED PLANNING
|   When: Safety-critical, high-risk operations
|   Approach: Formal specification, verification gates
|   Overhead: Very High
+-----+

```

### **Selection Decision Tree:**

```

Complexity Assessment
|
Is task atomic (1 step)?
+- YES -> DIRECT EXECUTION
|
+- NO -> Is path clear & sequential?
    +- YES -> LINEAR DECOMPOSITION
    |
    +- NO -> Is domain well-known?
        +- YES -> Do we have similar past tasks?
            |           +- YES -> CASE-BASED PLANNING
            |           |
            |           +- NO -> HTN PLANNING
            |
            +- NO -> Is it high-risk?
                +- YES -> VERIFIED PLANNING
                |
                +- NO -> EXPLORATORY PLANNING

```

## **2.4 Task Decomposer (HTN Focus)**

**HTN Planning Approach:** Inspired by classical HTN but adapted for LLM context.

### **Core Concepts:**

```

+-----+
|       HIERARCHICAL TASK NETWORK (HTN)
|
| Compound Task: High-level goal that needs decomposition
| Primitive Task: Atomic action directly executable
| Method: Way to decompose compound task into subtasks
| Ordering: Constraints on subtask execution sequence
+-----+

```

### **Decomposition Process:**

```

Goal: "Refactor authentication to use JWT"
|
+-----+
| Level 0: Goal (Compound Task)          |
| * Refactor authentication to JWT      |
+-----+
|
+-----+
| Level 1: Major Phases (Compound Tasks) |
| 1. Research & Design                 |
| 2. Implementation                     |
| 3. Testing & Validation              |
| 4. Deployment                        |
+-----+
|
+-----+
| Level 2: Concrete Steps (Mix)          |
| 1.1 Research JWT libraries (Compound) |
| 1.2 Design token flow (Compound)       |
| 2.1 Install JWT library (Primitive)    |
| 2.2 Implement token generation (Compound) |
| 2.3 Update auth middleware (Compound) |
| 3.1 Write unit tests (Compound)        |
| 3.2 Run test suite (Primitive)         |
| 4.1 Deploy to staging (Primitive)      |
| 4.2 Verify functionality (Compound)    |
+-----+
|
+-----+
| Level 3: Atomic Actions (Primitives)   |
| 1.1.1 Search for "jwt library python" |
| 1.1.2 Read library documentation       |
| 1.1.3 Compare options                  |
| 2.2.1 Create jwt_utils.py file         |
| 2.2.2 Import jwt library               |
| 2.2.3 Write encode_token function     |
| 2.2.4 Write decode_token function     |
| ... (continues)                      |
+-----+

```

### **Decomposition Algorithm (Conceptual):**

```

Function DECOMPOSE(task, depth):
  IF task is primitive:
    RETURN [task]

  IF depth > MAX_DEPTH:
    TREAT as primitive (avoid infinite recursion)
    RETURN [task]

```

```

# Query LLM for decomposition
methods = LLM.generate_methods(task, context)

# Select best method
method = SELECT_BEST(methods, criteria)

# Apply method to get subtasks
subtasks = APPLY_METHOD(method, task)

# Recursively decompose compound subtasks
plan = []
FOR subtask IN subtasks:
    plan.extend(DECOMPOSE(subtask, depth + 1))

RETURN plan

```

### **Method Selection Criteria:**

METHOD SELECTION FACTORS	
1. FEASIBILITY	<ul style="list-style-type: none"> <li>* Are required resources available?</li> <li>* Are dependencies satisfiable?</li> </ul>
2. EFFICIENCY	<ul style="list-style-type: none"> <li>* Estimated time</li> <li>* Estimated cost</li> <li>* Complexity overhead</li> </ul>
3. RISK	<ul style="list-style-type: none"> <li>* Probability of failure</li> <li>* Impact of failure</li> <li>* Reversibility</li> </ul>
4. QUALITY	<ul style="list-style-type: none"> <li>* Expected outcome quality</li> <li>* Maintainability</li> <li>* Robustness</li> </ul>
5. PAST PERFORMANCE	<ul style="list-style-type: none"> <li>* Success rate of similar methods</li> <li>* Learned preferences</li> </ul>

## **2.5 Dependency Analyzer**

**Purpose:** Costruire grafo delle dipendenze e identificare vincoli di ordinamento.

## **Dependency Types:**

DEPENDENCY TYPES	
1. PREREQUISITE (Must-Before)	Task B requires output of Task A Example: "Install library" before "Import library"
2. ORDERING (Should-Before)	Task B logically follows A (but not strict) Example: "Design" before "Implement"
3. RESOURCE (Shared-Resource)	Tasks compete for same resource, can't run parallel Example: Two tasks modifying same file
4. CONFLICT (Mutual-Exclusion)	Tasks logically conflict, must not overlap Example: "Backup DB" and "Modify DB"

## **Dependency Graph Construction:**

Subtasks: [T1, T2, T3, T4, T5]

+-----+	
Analyze Each Pair (Ti, Tj)	
* Does Ti produce input for Tj?	
* Do they access same resources?	
* Are they logically ordered?	
* Do they conflict?	
+-----+-----+	
+-----+	
Build Dependency Graph (DAG)	
T1	
/ \	
/ \	
T2    T3	
\ /	
\ /	
T4	
T5	
Edges represent dependencies	
+-----+-----+	

```

+
+-----+
| Identify Parallelizable Tasks      |
| * T2 and T3 can run in parallel   |
| (no dependencies between them)    |
+-----+
|
+-----+
| Compute Critical Path            |
| * Longest path through graph     |
| * Determines minimum execution time |
| * Example: T1 -> T2 -> T4 -> T5    |
+-----+

```

### **Topological Sort for Execution Order:**

Algorithm: Kahn's Algorithm (adapted)

Input: Dependency DAG

Output: Valid execution order

1. Compute in-degree for each task  
(number of dependencies)
2. Queue all tasks with in-degree 0  
(no dependencies)
3. While queue not empty:
  - a. Dequeue task T
  - b. Add T to execution order
  - c. For each task T' dependent on T:
    - Decrease in-degree of T'
    - If in-degree becomes 0, enqueue T'
4. If any tasks remain (not in order):
  - > Cycle detected! Error.
5. Return execution order

### **Example:**

Tasks: [Install\_JWT, Import\_JWT, Create\_Encoder, Create\_Decoder,  
Update\_Auth, Write\_Tests, Run\_Tests]

Dependencies:

```

Import_JWT -> Install_JWT
Create_Encoder -> Import_JWT
Create_Decoder -> Import_JWT
Update_Auth -> Create_Encoder, Create_Decoder
Write_Tests -> Update_Auth

```

Run\_Tests -> Write\_Tests

Execution Order (one valid ordering):

1. Install\_JWT
2. Import\_JWT
3. Create\_Encoder     } Can run in parallel
4. Create\_Decoder     }
5. Update\_Auth
6. Write\_Tests
7. Run\_Tests

Critical Path:

Install\_JWT -> Import\_JWT -> Create\_Encoder -> Update\_Auth  
-> Write\_Tests -> Run\_Tests

## 2.6 Resource Estimator

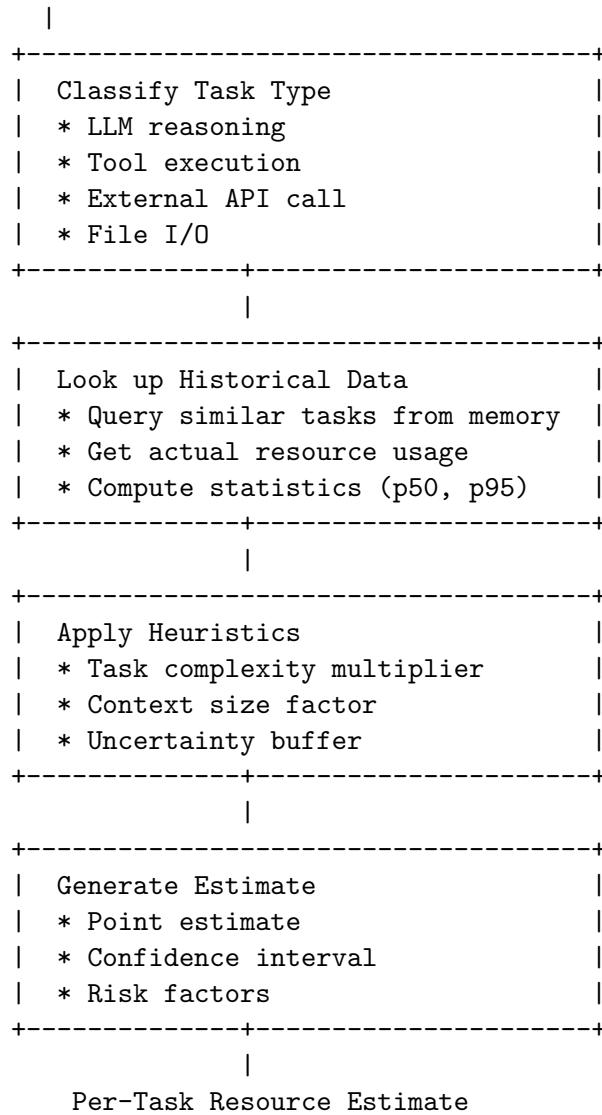
**Purpose:** Stimare risorse necessarie per ogni subtask e piano complessivo.

**Resource Types:**

RESOURCE DIMENSIONS	
1. TIME	
* Wall-clock time (latency)	
* LLM inference time	
* Tool execution time	
2. COMPUTE	
* LLM calls (count, size)	
* Token usage (input + output)	
* Model tier required (small/large)	
3. COST	
* LLM API cost	
* External API cost	
* Compute infrastructure cost	
4. MEMORY	
* Context window usage	
* Working memory size	
* Storage for intermediate results	
5. EXTERNAL DEPENDENCIES	
* API rate limits	
* Service availability	
* Network bandwidth	

## **Estimation Process:**

For each subtask:



## **Aggregation for Overall Plan:**

Individual Estimates -> Overall Plan Estimate

Sequential Tasks:

```
Time = sum(task_times)  
Cost = sum(task_costs)  
Tokens = sum(task_tokens)
```

Parallel Tasks:

```
Time = max(parallel_group_times)  
Cost = sum(parallel_group_costs)  
Tokens = sum(parallel_group_tokens)
```

With Uncertainty:

```
Time_estimate = expected_time * (1 + uncertainty_factor)
Cost_budget = expected_cost * (1 + buffer_factor)
```

### Example Estimation:

Task: "Refactor authentication to JWT"

Subtask Estimates:

Subtask	Time	LLM Calls	Tokens	Cost
Research JWT libs	30s	2	8K	\$0.02
Design token flow	45s	3	12K	\$0.03
Install library	10s	1	2K	\$0.01
Create encoder	60s	4	15K	\$0.04
Create decoder	60s	4	15K	\$0.04
Update auth	90s	5	20K	\$0.05
Write tests	120s	6	25K	\$0.06
Run tests	15s	0	0	\$0.00
<b>TOTAL (sequential)</b>	<b>430s</b>	<b>25</b>	<b>97K</b>	<b>\$0.25</b>
<b>TOTAL (optimized)</b>	<b>~300s</b>	<b>25</b>	<b>97K</b>	<b>\$0.25</b>
<b>(encoder+decoder in parallel)</b>				

Budget Check:

```
[v] Time: 300s < 600s budget
[v] Cost: $0.25 < $0.50 budget
[v] Tokens: 97K < 200K budget
-> Plan is feasible
```

## 2.7 Plan Optimizer

**Purpose:** Ottimizzare piano di esecuzione per obiettivi specificati.

### Optimization Objectives:

OPTIMIZATION OBJECTIVES	
1. MINIMIZE LATENCY	
* Maximize parallelization	
* Use faster models for non-critical tasks	
* Cache aggressively	
2. MINIMIZE COST	
* Use smaller models when possible	
* Batch operations	

```

|   * Reuse computations
|
| 3. MAXIMIZE QUALITY
|   * Use best models
|   * Add verification steps
|   * Multi-pass refinement
|
| 4. MINIMIZE RISK
|   * Add checkpoints
|   * Prefer reversible operations early
|   * Add approval gates
|
| 5. BALANCE (Default)
|   * Trade-off between above
|   * Context-appropriate weighting
+-----+

```

## **Optimization Techniques:**

### 1. PARALLELIZATION

Before: T1 → T2 → T3 → T4 (sequential)  
After: T1 → {T2, T3} → T4 (parallel)  
+-----+  
Speedup: 2x for T2+T3 phase

### 2. REORDERING

Before: Expensive → Cheap → Verify  
After: Cheap → Verify → Expensive  
Benefit: Fail fast, save expensive work

### 3. BATCHING

Before: API\_call(x1), API\_call(x2), API\_call(x3)  
After: API\_batch\_call([x1, x2, x3])  
Benefit: Reduce overhead, cost

### 4. CACHING

Before: Compute F(x) multiple times  
After: Compute once, reuse result  
Benefit: Eliminate redundant work

### 5. MODEL ROUTING

Before: Use large model for all tasks  
After: Use small model for simple, large for complex  
Benefit: Reduce cost without quality loss

### 6. INCREMENTAL EXECUTION

Before: Build entire solution, then test  
After: Build incrementally, test each piece  
Benefit: Earlier feedback, easier debugging

## **Optimization Algorithm** (Conceptual):

```
Function OPTIMIZE_PLAN(plan, objective):  
  
    # 1. Identify parallelization opportunities  
    parallel_groups = FIND_PARALLEL_TASKS(plan.dag)  
    plan = APPLY_PARALLELIZATION(plan, parallel_groups)  
  
    # 2. Reorder for fail-fast  
    IF objective prioritizes risk_minimization:  
        plan = REORDER_FOR_EARLY_VALIDATION(plan)  
  
    # 3. Apply caching  
    cacheable = IDENTIFY_CACHEABLE_COMPUTATIONS(plan)  
    plan = INSERT_CACHE_LOOKUPS(plan, cacheable)  
  
    # 4. Model routing optimization  
    FOR task IN plan.tasks:  
        IF task is simple AND objective prioritizes cost:  
            task.model = SMALL_MODEL  
        ELSE IF task is critical AND objective prioritizes quality:  
            task.model = LARGE_MODEL  
        ELSE:  
            task.model = MEDIUM_MODEL  
  
    # 5. Batch operations  
    batchable = FIND_BATCHABLE_OPERATIONS(plan)  
    plan = APPLY_BATCHING(plan, batchable)  
  
    RETURN plan
```

## **2.8 Contingency Planner**

**Purpose:** Preparare strategie di fallback per gestire failure previsti.

### **Failure Taxonomy:**

FAILURE TYPES	
1. TRANSIENT FAILURES	
* Network timeout	
* Rate limit hit	
* Temporary service unavailable	
Recovery: Retry with backoff	
2. RESOURCE FAILURES	
* Budget exhausted	
* Timeout exceeded	

```

|   * Memory limit reached
|   Recovery: Abort gracefully, report
|
| 3. LOGIC FAILURES
|   * Task prerequisites not met
|   * Invalid input to tool
|   * Assertion violation
|   Recovery: Replan, adjust approach
|
| 4. EXTERNAL FAILURES
|   * Tool unavailable
|   * API breaking change
|   * File not found
|   Recovery: Use alternative tool/approach
|
| 5. QUALITY FAILURES
|   * Output doesn't meet criteria
|   * Test failures
|   * Validation rejection
|   Recovery: Refine, iterate, or escalate
+-----+

```

### **Contingency Planning Process:**

For each subtask in plan:

```

| 
+-----+
| Identify Potential Failures
|   * What can go wrong?
|   * Historical failure patterns
|   * Uncertainty indicators
+-----+
| 
+-----+
| Assess Failure Impact
|   * Blocks other tasks?
|   * Partial vs complete failure
|   * Reversibility of side effects
+-----+
| 
+-----+
| Generate Recovery Strategies
|   * Retry with modifications
|   * Alternative approaches
|   * Fallback to simpler solution
|   * Human escalation
+-----+
| 
+-----+

```

```

| Attach to Plan           |
| * Each task has contingency list   |
| * Ordered by preference      |
| * Conditions for triggering    |
+-----+

```

## **Contingency Structure:**

```

Contingency {
    failure_type: FailureType,
    detection: {
        symptoms: [observable_conditions],
        confidence_threshold: float
    },
    recovery_strategies: [
        {
            strategy_id: string,
            description: string,
            preconditions: [conditions_to_apply],
            actions: [recovery_steps],
            expected_success_rate: float,
            cost: ResourceEstimate
        }
    ],
    escalation: {
        max_retries: int,
        escalate_to: "human" | "alternative_plan",
        escalation_message: string
    }
}

```

## **Example Contingencies:**

Task: "Install JWT library"

Contingency 1:

Failure: "Package not found in pip"

Recovery Strategies:

1. Try alternative package name (e.g., PyJWT vs python-jwt)
2. Search PyPI for similar packages
3. Install from source (GitHub)
4. Ask user for guidance

Contingency 2:

Failure: "Permission denied"

Recovery Strategies:

1. Retry with sudo (if appropriate context)
2. Install in virtual environment
3. Install with --user flag
4. Request user to fix permissions

Contingency 3:

- Failure: "Network timeout"
- Recovery Strategies:
  1. Retry with exponential backoff (3 attempts)
  2. Try alternative package index
  3. Defer to later (maybe network recovers)
  4. Report failure, request manual installation

## 2.9 Planning Engine Output Schema

```
ExecutionPlan {
  // Metadata
  plan_id: string,
  created_at: datetime,
  goal_analysis_id: string,
  strategy_used: PlanningStrategy,

  // Task Structure
  tasks: [
    {
      task_id: string,
      description: string,
      type: "primitive" | "compound",
      parent_task_id: string | null,
      depth: int,

      // Execution Details
      action: {
        type: "llm_call" | "tool_invocation" | "composite",
        parameters: {...},
        model: ModelSpec,
        tools: [ToolSpec]
      },
    },
    ...
  ],
  // Dependencies
  prerequisites: [task_id],
  produces: [output_id],
  consumes: [input_id],

  // Resource Estimates
  estimates: {
    time: {expected: float, confidence_interval: [float, float]},
    cost: {expected: float, max: float},
    tokens: {input: int, output: int}
  },
}

// Contingencies
```

```

        contingencies: [Contingency],
        // Verification
        success_criteria: [Criterion],
        verification_method: string
    }
],
// Dependency Structure
dependency_graph: {
    nodes: [task_id],
    edges: [{from: task_id, to: task_id, type: DependencyType}]
},
// Execution Order
execution_sequence: [
    {
        phase: int,
        parallel_groups: [[task_id]] // Tasks in same group can run in parallel
    }
],
// Overall Estimates
overall_estimates: {
    min_time: float,
    expected_time: float,
    max_time: float,
    expected_cost: float,
    max_cost: float,
    total_tokens: int
},
// Optimization Info
optimization: {
    objective: OptimizationObjective,
    parallelization_opportunities: int,
    cached_computations: int
},
// Risk Assessment
risk: {
    overall_risk: RiskLevel,
    high_risk_tasks: [task_id],
    failure_probability: float,
    mitigation_strategies: [string]
}
}

```

### 3. Execution Engine

#### 3.1 Purpose & Responsibilities

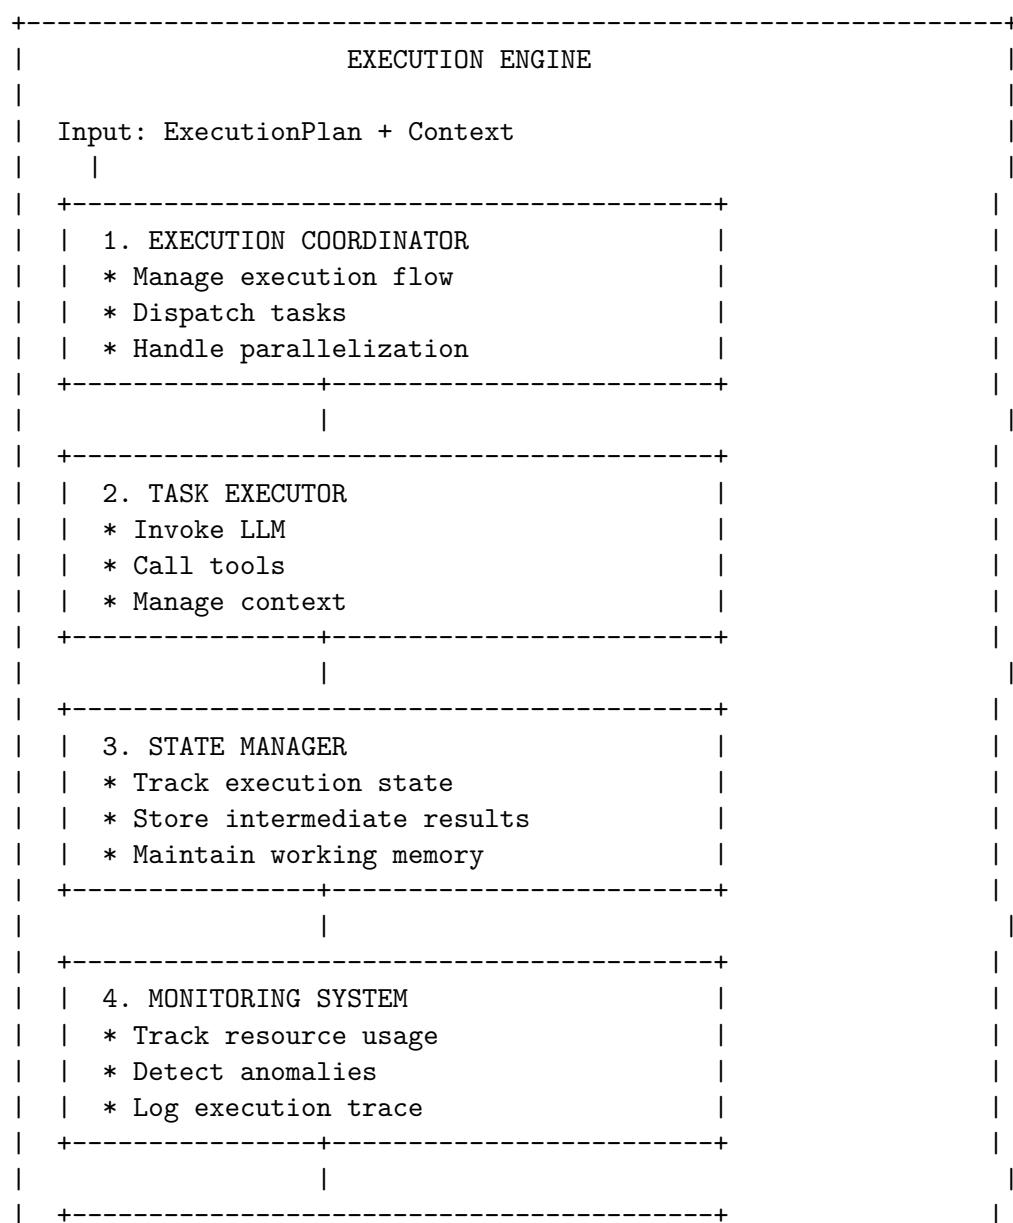
**Core Function:** Eseguire il plan generato, monitorando progresso, verificando risultati, e adattando dinamicamente quando necessario.

**Responsibilities:** 1. **Task Execution:** Invocare azioni (LLM calls, tool usage) 2.

**State Management:** Tracciare stato corrente, output intermedi 3. **Monitoring:** Osservare progresso, resource usage, problemi 4. **Verification:** Validare output contro success criteria 5. **Adaptation:** Replan o adjust quando si incontrano ostacoli 6.

**Safety:** Enforce bounds, prevent unsafe actions

#### 3.2 Architecture Interna



```

| | 5. VERIFICATION LAYER |
| | * Validate outputs |
| | * Check success criteria |
| | * Safety verification |
| +-----+
| |
| +-----+
| | 6. ADAPTATION ENGINE |
| | * Detect need for replanning |
| | * Apply contingencies |
| | * Dynamic strategy adjustment |
| +-----+
| |
| Output: ExecutionResult
| {
|   status: SUCCESS | FAILURE | PARTIAL,
|   outputs: {...},
|   trace: [...],
|   stats: {...}
| }
+-----+

```

### 3.3 Execution Coordinator

**Purpose:** Orchestrare esecuzione di piano, gestendo dipendenze e parallelizzazione.

**Execution Loop** (Main Algorithm):

Function EXECUTE\_PLAN(plan):

```

# Initialize
state = INITIALIZE_STATE(plan)
execution_trace = []

# Main execution loop
FOR phase IN plan.execution_sequence:

    # Execute parallel groups in this phase
    FOR parallel_group IN phase.parallel_groups:

        # Check if all prerequisites satisfied
        IF NOT ALL_PREREQUISITES_MET(parallel_group, state):
            HANDLE_BLOCKED_EXECUTION(parallel_group, state)
            CONTINUE

        # Execute tasks in parallel (if multiple)
        IF len(parallel_group) > 1:
            results = PARALLEL_EXECUTE(parallel_group, state)
        ELSE:

```

```

results = [EXECUTE_TASK(parallel_group[0], state)]

# Process results
FOR task, result IN zip(parallel_group, results):

    # Verify result
    verification = VERIFY_RESULT(task, result, state)

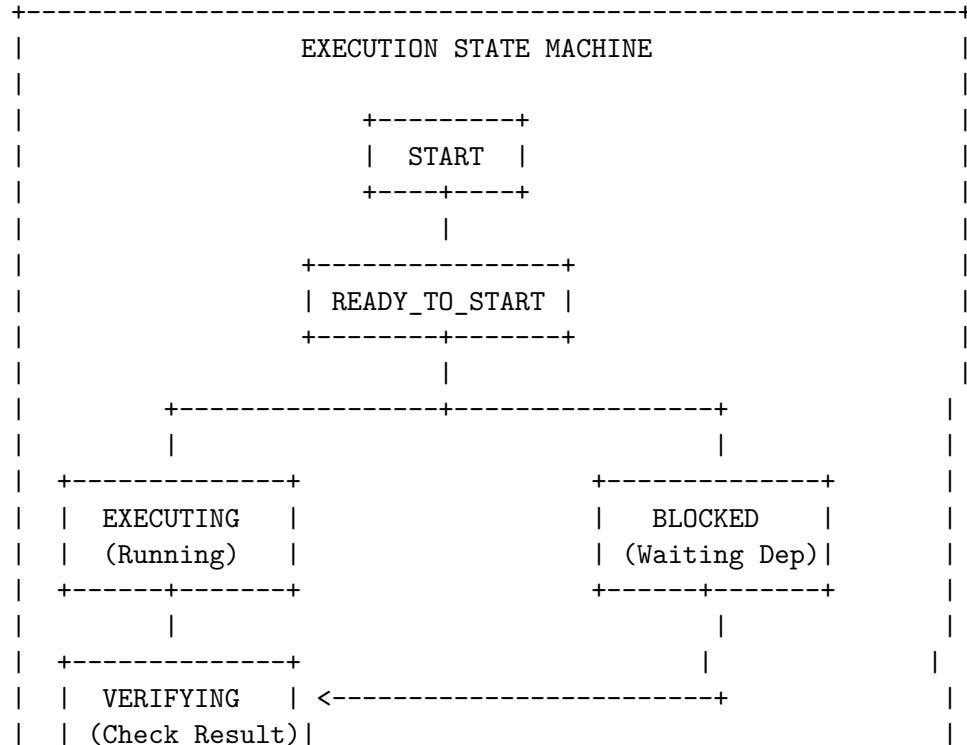
    IF verification.success:
        # Update state with successful result
        state = UPDATE_STATE(state, task, result)
        execution_trace.append({task, result, "SUCCESS"})
    ELSE:
        # Handle failure
        recovery = APPLY_CONTINGENCY(task, result, verification)

        IF recovery.success:
            state = UPDATE_STATE(state, task, recovery.result)
            execution_trace.append({task, recovery.result, "RECOVERED"})
        ELSE:
            # Failure not recoverable at task level
            RETURN PARTIAL_RESULT(state, execution_trace, task)

    # All phases completed
RETURN SUCCESS_RESULT(state, execution_trace)

```

### **State Machine:**



```

+-----+
| |
+-----+
| | Verification Pass? |
| +-----+---+
| | YES | NO |
| | | |
+-----+ +-----+
| | SUCCESS | | FAILED |
| +-----+ +-----+
| |
+-----+
| | Contingency OK? |
| +-----+---+
| | YES | NO |
| | | |
+-----+ +-----+
| | RECOVERING | | ABORTED |
| +-----+ +-----+
| |
| (Back to EXECUTING)
+-----+

```

### 3.4 Task Executor

**Purpose:** Eseguire singoli task (primitive actions).

**Execution Types:**

```

+-----+
| |
+-----+
| | TASK EXECUTION TYPES |
| |
| | 1. LLM_CALL
| |   * Prepare context (system + user prompt)
| |   * Route to appropriate model
| |   * Invoke with tools if needed
| |   * Extract structured output
| |
| | 2. TOOL_INVOCATION
| |   * Retrieve tool from registry
| |   * Validate inputs
| |   * Execute tool with safety bounds
| |   * Capture output
| |
| | 3. COMPOUND_ACTION
| |   * Execute sub-plan recursively
| |   * Aggregate results
| |
| | 4. VERIFICATION_STEP
| |
+-----+

```

```
| * Run validation logic |  
| * Check assertions |  
| * Produce verification report |  
+-----+
```

## LLM Call Execution Flow:

Task: "Generate JWT token encoding function"

```
|  
+-----+  
| 1. Context Preparation |  
|  
| System Prompt:  
| "You are expert Python developer"  
|  
| User Prompt:  
| "Generate JWT encoding function  
| Requirements: [...]"  
|  
| Working Memory:  
| - JWT library choice (PyJWT)  
| - Previous implementation context  
+-----+-----+  
|  
+-----+  
| 2. Model Routing |  
| * Task complexity: MEDIUM |  
| * Context size: 5K tokens |  
| * Budget remaining: OK |  
| -> Route to: claude-sonnet-3.5 |  
+-----+-----+  
|  
+-----+  
| 3. LLM Invocation |  
| * Model: claude-sonnet-3.5 |  
| * Max tokens: 2000 |  
| * Temperature: 0.2 |  
| * Tools: [write_file, read_file] |  
+-----+-----+  
|  
+-----+  
| 4. Output Extraction |  
| * Parse LLM response |  
| * Extract function code |  
| * Extract any metadata |  
+-----+-----+  
|  
Task Result
```

### Tool Invocation Flow:

```
Tool: "write_file"
Parameters: {path: "jwt_utils.py", content: "..."}
|
+-----+
| 1. Safety Verification      |
| * Check path against whitelist |
| * Verify not overwriting critical |
| * Check size limits          |
+-----+
|
+-----+
| 2. Input Validation        |
| * Schema validation        |
| * Type checking            |
| * Constraint verification   |
+-----+
|
+-----+
| 3. Tool Execution          |
| * Invoke tool with parameters |
| * Capture stdout/stderr     |
| * Monitor for errors       |
+-----+
|
+-----+
| 4. Output Capture          |
| * Parse tool output         |
| * Structure as ToolResult  |
| * Attach metadata           |
+-----+
|
Tool Result
```

### 3.5 State Manager

**Purpose:** Mantenere stato corrente di esecuzione e risultati intermedi.

#### State Structure:

```
ExecutionState {
    // Task Status
    completed_tasks: {task_id: Result},
    in_progress_tasks: {task_id: StartTime},
    pending_tasks: [task_id],
    failed_tasks: {task_id: FailureInfo},

    // Working Memory (context for execution)
```

```

working_memory: {
    task_outputs: {output_id: Value},
    intermediate_results: {key: Value},
    context_variables: {var_name: Value}
},
// Resource Tracking
resources_used: {
    time_elapsed: float,
    tokens_consumed: int,
    cost_accumulated: float,
    llm_calls: int,
    tool_invocations: int
},
// Dependencies
dependency_satisfaction: {
    task_id: boolean // Are all prerequisites met?
},
// Errors & Warnings
errors: [Error],
warnings: [Warning],
// Metadata
execution_start: datetime,
last_update: datetime
}

```

## **State Transitions:**

State Update Operations:

1. TASK\_STARTED(task\_id, timestamp):
 

```

pending_tasks.remove(task_id)
in_progress_tasks[task_id] = timestamp
      
```
2. TASK\_COMPLETED(task\_id, result):
 

```

in_progress_tasks.remove(task_id)
completed_tasks[task_id] = result
working_memory.task_outputs[task_id] = result.output
UPDATE_DEPENDENCY_SATISFACTION(task_id)
      
```
3. TASK\_FAILED(task\_id, error):
 

```

in_progress_tasks.remove(task_id)
failed_tasks[task_id] = error
errors.append(error)
      
```
4. UPDATE\_RESOURCES(resource\_delta):
 

```

      
```

```

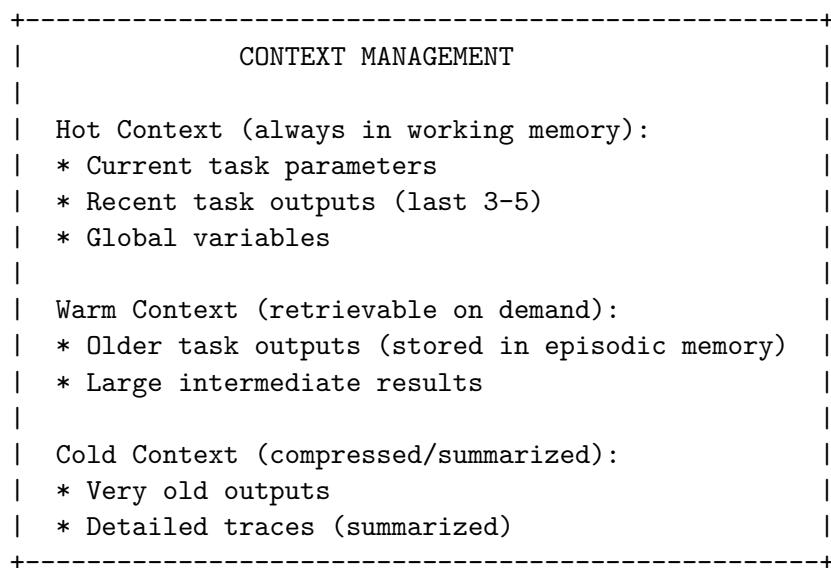
    resources_used += resource_delta

5. SET_CONTEXT(key, value):
    working_memory.context_variables[key] = value

```

### **Context Management:**

Working Memory has limited size -> Must manage carefully



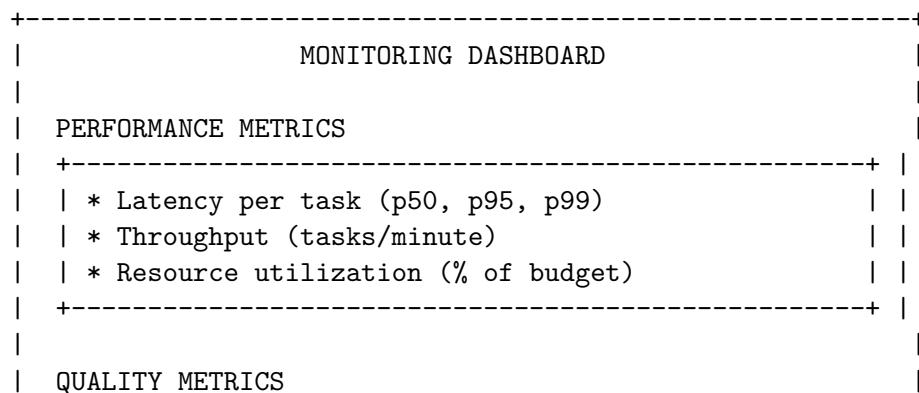
Context Size Budget: ~20K tokens

- If approaching limit:
  1. Summarize old outputs
  2. Move to episodic memory
  3. Keep only essential info

## **3.6 Monitoring System**

**Purpose:** Osservare esecuzione in real-time, detect anomalies, collect telemetry.

### **Monitoring Dimensions:**



```

| +-----+ |
| | * Task success rate | |
| | * Verification pass rate | |
| | * Error frequency | |
| +-----+ |
|
| ANOMALY DETECTION |
| +-----+ |
| | * Task taking unusually long | |
| | * Repeated failures on same task | |
| | * Resource usage spike | |
| | * Unexpected error patterns | |
| +-----+ |
|
| EXECUTION TRACE |
| +-----+ |
| | * Detailed log of all actions | |
| | * Inputs/outputs for each task | |
| | * Decision points and reasoning | |
| | * Timestamps and durations | |
| +-----+ |
+-----+

```

### **Anomaly Detection:**

Monitor continuously for:

1. LATENCY ANOMALY  
 IF task\_duration > expected\_duration \* 3:  
   WARN: "Task taking unusually long"  
   CONSIDER: Timeout, interrupt, or investigate
2. REPEATED FAILURE  
 IF same\_task\_failed > 3 times:  
   WARN: "Repeated failures detected"  
   CONSIDER: Skip task, try alternative, escalate
3. RESOURCE SPIKE  
 IF resource\_rate > budget\_rate \* 2:  
   WARN: "Resource consumption too high"  
   CONSIDER: Throttle, optimize, or abort
4. ERROR PATTERN  
 IF similar\_errors > threshold:  
   WARN: "Systematic error detected"  
   CONSIDER: Diagnose root cause, replan
5. STUCK EXECUTION  
 IF no\_progress\_for > timeout:

```
WARN: "Execution appears stuck"
CONSIDER: Interrupt, diagnose
```

## Logging:

```
Execution Trace Entry {
    timestamp: datetime,
    task_id: string,
    event_type: "START" | "COMPLETE" | "FAIL" | "VERIFY",

    // Context
    inputs: {...},
    working_memory_snapshot: {...},

    // Execution Details
    action_taken: string,
    llm_prompt: string (if applicable),
    llm_response: string (if applicable),
    tool_calls: [ToolCall],

    // Results
    output: {...},
    verification_result: {...},

    // Metadata
    duration: float,
    tokens_used: int,
    cost: float,
    model_used: string,

    // Decisions
    decisions_made: [Decision],
    reasoning: string
}
```

## 3.7 Verification Layer

**Purpose:** Validare output di ogni task contro success criteria e safety bounds.

### Verification Types:

VERIFICATION HIERARCHY	
Level 1: SCHEMA VERIFICATION	
* Output structure matches expected format	
* Required fields present	
* Type correctness	
-> Fast, automatic, always run	

```

| Level 2: CONSTRAINT VERIFICATION
| * Output satisfies explicit constraints
| * Bounds checking (numeric ranges, lengths)
| * Format validation (regex patterns)
| -> Fast, automatic, always run
|
| Level 3: SEMANTIC VERIFICATION
| * Output semantically correct
| * Meaningful in context
| * Logically consistent
| -> Moderate cost, LLM-based, selective
|
| Level 4: FUNCTIONAL VERIFICATION
| * Output achieves intended goal
| * Side effects as expected
| * Integration with system works
| -> Expensive, testing required, critical tasks
|
| Level 5: SAFETY VERIFICATION
| * No prohibited actions taken
| * No security vulnerabilities introduced
| * Compliance requirements met
| -> Critical, always run for sensitive operations
+-----+

```

## **Verification Flow:**

Task Result

```

| |
+-----+
| 1. Schema Verification
| * Check structure
| * Validate types
| IF fail: REJECT immediately
+-----+
| PASS
+-----+
| 2. Constraint Verification
| * Check bounds
| * Validate format
| IF fail: REJECT or REQUEST_RETRY
+-----+
| PASS
+-----+
| 3. Safety Verification
| * Scan for prohibited patterns
| * Check against safety rules
| IF fail: REJECT + ALERT
+-----+

```

```

+-----+
| PASS
+-----+
| 4. Semantic Verification (Optional) |
| * LLM-based quality check          |
| * Context appropriateness         |
| IF fail: FLAG for review        |
+-----+
| PASS
+-----+
| 5. Functional Verification (Select) |
| * Execute tests                   |
| * Verify side effects            |
| IF fail: RETRY with fixes       |
+-----+
| PASS
ACCEPT Result

```

### **Success Criteria Evaluation:**

Task has success\_criteria: [Criterion]

```

Criterion {
    type: "OUTPUT_PROPERTY" | "SIDE_EFFECT" | "TEST_PASS" | "CUSTOM",
    description: string,
    verification_method: {
        type: "AUTOMATIC" | "LLM_CHECK" | "TOOL_EXECUTION",
        parameters: {...}
    },
    required: boolean // Must pass vs nice-to-have
}

```

Evaluation Process:

```

FOR criterion IN task.success_criteria:
    result = VERIFY_CRITERION(criterion, task_output)

    IF criterion.required AND NOT result.passed:
        RETURN VERIFICATION_FAILED(criterion, result.reason)

    IF NOT criterion.required AND NOT result.passed:
        LOG_WARNING(criterion, result.reason)

RETURN VERIFICATION_PASSED

```

### **Example Verification:**

Task: "Write JWT encoding function"  
Output: (Python function code)

Success Criteria:

1. Function is valid Python (REQUIRED)  
Method: Parse with ast module  
Result: [v] PASS
2. Function accepts correct parameters (REQUIRED)  
Method: Check signature: encode\_token(payload, secret)  
Result: [v] PASS
3. Function returns string (REQUIRED)  
Method: Check return type hint  
Result: [v] PASS
4. Function uses PyJWT library (REQUIRED)  
Method: Check for "import jwt" and "jwt.encode"  
Result: [v] PASS
5. Function handles errors gracefully (OPTIONAL)  
Method: Check for try-except block  
Result: [v] PASS
6. Function has docstring (OPTIONAL)  
Method: Check for docstring  
Result: [x] FAIL (but optional, so warning only)

Overall: PASS (all required criteria met)

### 3.8 Adaptation Engine

**Purpose:** Detect quando execution devia da plan e adattarsi dinamicamente.

**Adaptation Triggers:**

ADAPTATION TRIGGERS	
1. TASK FAILURE	
* Task execution failed	
* Verification rejected output	
-> Apply contingency or replan	
2. ASSUMPTION VIOLATION	
* Expected precondition not met	
* Resource not available	
-> Update assumptions, replan affected subtasks	
3. OPPORTUNITY DISCOVERY	
* Better approach discovered during execution	
* Shortcut found	
-> Optimize plan opportunistically	

4. RESOURCE PRESSURE	
* Budget running low	
* Time pressure increased	
-> Switch to more efficient strategy	
5. CONTEXT CHANGE	
* External environment changed	
* Requirements updated	
-> Replan with new context	

## Adaptation Strategies:

	ADAPTATION STRATEGIES
	STRATEGY 1: CONTINGENCY APPLICATION
	When: Anticipated failure
	Action: Apply pre-planned fallback
	Cost: Low (already planned)
	Example: Package not found -> Try alternative name
	STRATEGY 2: LOCAL REPLANNING
	When: Unanticipated failure, affects single task
	Action: Replan just the failed task
	Cost: Medium (limited replanning)
	Example: API changed -> Find new way to call it
	STRATEGY 3: SUBTREE REPLANNING
	When: Failure affects multiple dependent tasks
	Action: Replan failed task and all descendants
	Cost: Medium-High (more extensive)
	Example: Library choice wrong -> Redo all library-specific
	STRATEGY 4: GLOBAL REPLANNING
	When: Fundamental assumption violated
	Action: Restart planning from current state
	Cost: High (redo most work)
	Example: Requirements completely changed
	STRATEGY 5: ESCALATION
	When: Adaptation not possible within bounds
	Action: Request human intervention
	Cost: Blocks progress until human responds
	Example: Ambiguity can't be resolved automatically

## Adaptation Decision Tree:

```

Failure Detected
|
Has pre-planned contingency?
+- YES -> Apply Contingency
|   |
|   Success?
|     +- YES -> Continue Execution
|     +- NO -> Proceed to replanning
|
+- NO -> Is failure isolated?
  +- YES (single task) -> Local Replan
    |           |
    |           Replan one task
    |           Continue execution
    |
  +- NO (affects multiple) -> Dependency Analysis
    |
    How many tasks affected?
      +- Few (<5) -> Subtree Replan
        |           |
        |           Replan affected subtree
        |
      +- Many (>5) -> Global Replan?
        |
        Check if viable
          +- YES -> Global Replan
            |           (restart from current)
            |
          +- NO -> Escalate to Human

```

### **Adaptation Algorithm:**

```

Function ADAPT(failure, state, plan):

  # Classify failure severity
  severity = CLASSIFY_FAILURE(failure)

  # Check for contingency
  IF plan.has_contingency_for(failure):
    contingency = plan.get_contingency(failure)
    result = APPLY_CONTINGENCY(contingency, state)

    IF result.success:
      RETURN CONTINUE_EXECUTION(state)

  # Analyze impact
  affected_tasks = FIND_AFFECTED_TASKS(failure, plan)

  # Select adaptation strategy

```

```

IF len(affected_tasks) == 1:
    # Local replanning
    new_task_plan = REPLAN_SINGLE_TASK(affected_tasks[0], state)
    plan = REPLACE_TASK(plan, affected_tasks[0], new_task_plan)
    RETURN CONTINUE_EXECUTION(state)

ELSE IF len(affected_tasks) <= 5:
    # Subtree replanning
    new_subtree_plan = REPLAN_SUBTREE(affected_tasks, state)
    plan = REPLACE_SUBTREE(plan, affected_tasks, new_subtree_plan)
    RETURN CONTINUE_EXECUTION(state)

ELSE IF severity < CRITICAL AND resources_allow_replan:
    # Global replanning
    new_plan = REPLAN_FROM_SCRATCH(original_goals, state)
    RETURN RESTART_EXECUTION(new_plan, state)

ELSE:
    # Escalate
    RETURN ESCALATE_TO_HUMAN(failure, state, reason)

```

### **Example Adaptation:**

Scenario: Installing JWT library

Original Plan:

1. Run: pip install PyJWT
2. Import jwt in code
3. Use jwt.encode()

Execution:

Step 1: pip install PyJWT  
-> FAILURE: "Package not found"

Adaptation:

1. Check Contingency:  
[v] Found: "Try alternative package names"

2. Apply Contingency:  
Attempt 1: pip install python-jwt  
-> FAILURE: "Package not found"

Attempt 2: Search PyPI for "jwt python"  
-> SUCCESS: Found packages [PyJWT, python-jose, authlib]

Attempt 3: pip install python-jose  
-> SUCCESS: Installed

3. Local Replan:  
Affected task: Step 2 (Import statement)  
Old: import jwt  
New: from jose import jwt

Affected task: Step 3 (API usage)  
Old: jwt.encode(...)  
New: jwt.encode(...) # (API compatible)

4. Continue Execution:  
Execute updated steps 2 and 3  
-> SUCCESS

Outcome: Task completed successfully via adaptation

### 3.9 Execution Engine Output Schema

```
ExecutionResult {  
    // Status  
    status: "SUCCESS" | "PARTIAL_SUCCESS" | "FAILURE",  
    completion_percentage: float, // 0-100%  
  
    // Outputs  
    primary_output: {...}, // Main result  
    intermediate_outputs: {task_id: output},  
    side_effects: [SideEffect],  
  
    // Execution Trace  
    trace: [  
        {  
            task_id: string,  
            timestamp: datetime,  
            event: "START" | "COMPLETE" | "FAIL" | "ADAPT",  
            details: {...},  
            duration: float,  
            resources_used: ResourceDelta  
        }  
    ],  
  
    // Task Results  
    completed_tasks: [  
        {  
            task_id: string,  
            result: {...},  
            verification: VerificationResult,  
            duration: float  
        }  
    ],
```

```

failed_tasks: [
  {
    task_id: string,
    error: Error,
    attempted_contingencies: [Contingency],
    reason_for_failure: string
  }
],
// Adaptations
adaptations_applied: [
  {
    trigger: AdaptationTrigger,
    strategy: AdaptationStrategy,
    tasks_affected: [task_id],
    outcome: "SUCCESS" | "FAILURE"
  }
],
// Resource Usage
resources: {
  total_time: float,
  llm_calls: int,
  tokens_consumed: {input: int, output: int},
  cost: float,
  tool_invocations: {tool_name: count}
},
// Quality Metrics
metrics: {
  plan_adherence: float, // % of original plan executed
  adaptation_frequency: float, // # adaptations / # tasks
  verification_pass_rate: float, // % tasks passed verification
  efficiency: float // actual_time / estimated_time
},
// Errors & Warnings
errors: [Error],
warnings: [Warning],
// Context for Reflection
learning_insights: [
  {
    insight: string,
    category: "STRATEGY" | "FAILURE_PATTERN" | "OPTIMIZATION",
    confidence: float
  }
],

```

```

    // Metadata
    execution_id: string,
    plan_id: string,
    start_time: datetime,
    end_time: datetime
}

```

## 4. Reflection Module

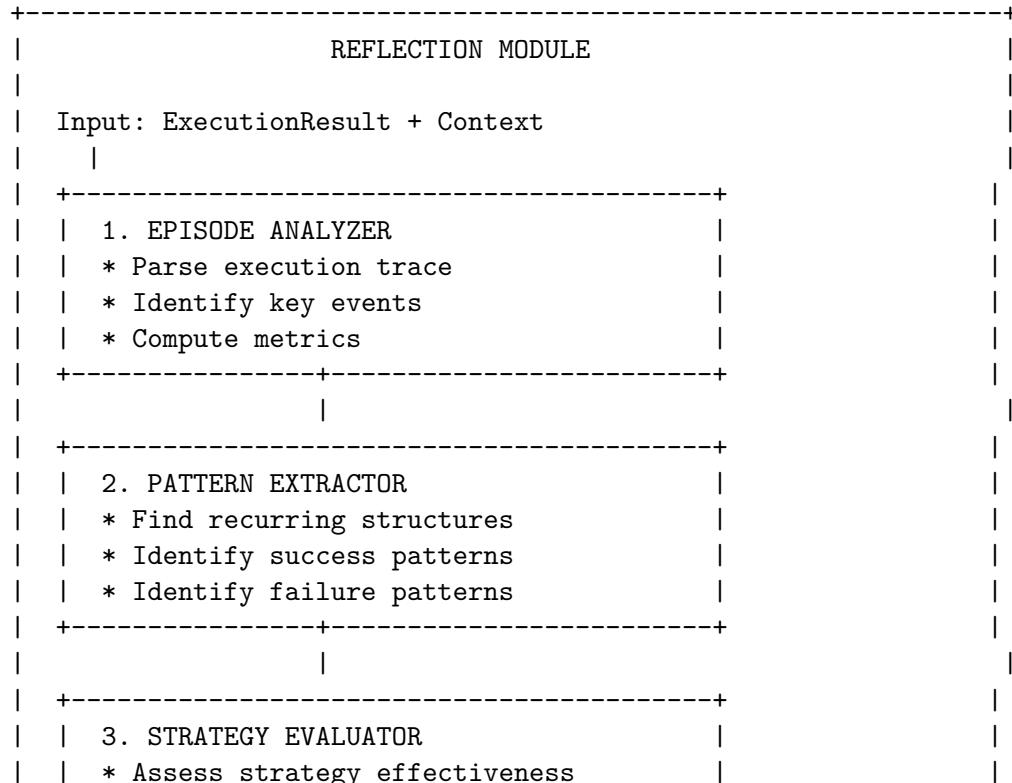
### 4.1 Purpose & Responsibilities

**Core Function:** Analizzare episodi di esecuzione completati per estrarre learning, identificare pattern, e migliorare performance future.

**Key Distinction:** Reflection è **post-execution**, non real-time. Avviene DOPO task completion per informare future executions.

**Responsibilities:** 1. **Episode Analysis:** Analizzare execution trace per successi/fallimenti 2. **Pattern Extraction:** Identificare pattern ricorrenti 3. **Strategy Learning:** Capire quali strategie funzionano in quali contesti 4. **Performance Analysis:** Valutare efficiency vs expectations 5. **Knowledge Distillation:** Convertire esperienze in reusable knowledge 6. **Memory Update:** Aggiornare episodic memory e pattern cache

### 4.2 Architecture Interna



```

|   |   * Compare alternatives
|   |   * Update strategy preferences
|   +-----+
|   |
|   +-----+
|   |   4. PERFORMANCE ANALYZER
|   |   * Actual vs estimated comparison
|   |   * Bottleneck identification
|   |   * Efficiency opportunities
|   +-----+
|   |
|   +-----+
|   |   5. KNOWLEDGE DISTILLER
|   |   * Generalize from specific
|   |   * Create reusable heuristics
|   |   * Update beliefs
|   +-----+
|   |
|   +-----+
|   |   6. MEMORY UPDATER
|   |   * Store episode in episodic memory
|   |   * Update pattern cache
|   |   * Consolidate knowledge
|   +-----+
|   |
|   Output: ReflectionInsights + Memory Updates
+-----+

```

### 4.3 Episode Analyzer

**Purpose:** Processare execution trace per estrarre eventi chiave e metriche.

**Analysis Dimensions:**

```

+-----+
|           EPISODE ANALYSIS
|
|           TEMPORAL ANALYSIS
|   * Total duration
|   * Time per phase
|   * Time per task
|   * Critical path actual time
|
|           SUCCESS/FAILURE ANALYSIS
|   * Which tasks succeeded
|   * Which tasks failed
|   * Failure reasons taxonomy
|   * Recovery success rate
|

```

RESOURCE ANALYSIS	
* Tokens consumed vs estimated	
* Cost actual vs budget	
* Tool usage patterns	
* Model usage patterns	
QUALITY ANALYSIS	
* Verification pass rates	
* Output quality scores	
* Rework frequency	
ADAPTATION ANALYSIS	
* # of adaptations	
* Adaptation triggers	
* Adaptation success rate	
* Plan adherence percentage	

### **Key Events Identification:**

Execution Trace → Key Events

Event Types to Extract:

1. DECISION POINTS
  - \* Where agent chose between alternatives
  - \* Reasoning for choice
  - \* Outcome of choice
2. FAILURES
  - \* What failed
  - \* Why it failed
  - \* How it was (or wasn't) recovered
3. ADAPTATIONS
  - \* What triggered adaptation
  - \* How agent adapted
  - \* Effectiveness of adaptation
4. BOTTLENECKS
  - \* Tasks that took unexpectedly long
  - \* Resource constraints hit
  - \* Blocking dependencies
5. SURPRISES
  - \* Unexpected successes
  - \* Unexpected failures
  - \* Assumption violations

## **Metrics Computation:**

Computed Metrics:

### EFFICIENCY METRICS:

```
efficiency_ratio = actual_time / estimated_time  
cost_efficiency = actual_cost / estimated_cost  
resource_utilization = resources_used / resources_available
```

### QUALITY METRICS:

```
success_rate = completed_tasks / total_tasks  
first_time_success = tasks_passed_verification_first_try / total_tasks  
rework_rate = tasks_requiring_retry / total_tasks
```

### ADAPTATION METRICS:

```
adaptation_rate = adaptations / total_tasks  
adaptation_success = successful_adaptations / total_adaptations  
plan_stability = tasks_from_original_plan / total_tasks_executed
```

### LEARNING METRICS:

```
novel_situations = situations_not_in_memory / total_situations  
pattern_reuse = patterns_applied_from_cache / opportunities
```

## **4.4 Pattern Extractor**

**Purpose:** Identificare pattern ricorrenti che possono informare future decisioni.

### **Pattern Types:**

PATTERN TAXONOMY	
1. SUCCESS PATTERNS	
* Task sequences that work well	
* Effective strategies for problem types	
* Optimal tool combinations	
Example: "For Python refactoring, always run tests before and after"	
2. FAILURE PATTERNS	
* Common failure modes	
* Error co-occurrences	
* Problematic assumptions	
Example: "Installing package X often fails on first try, but succeeds after pip upgrade"	
3. OPTIMIZATION PATTERNS	
* Tasks that can be parallelized	
* Caching opportunities	

```

|   * Shortcuts and simplifications      |
|   Example: "Can cache library search results for 1hr"  |
|
| 4. CONTEXT PATTERNS                  |
|   * Problem characteristics -> best approach          |
|   * Environment signals -> strategy selection          |
|   Example: "If codebase is TypeScript, verification      |
|             must include type checking"                   |
|
| 5. TEMPORAL PATTERNS                |
|   * Time-of-day effects              |
|   * Sequence effects (order matters) |
|   Example: "Database operations faster in mornings"    |
+-----+

```

### **Pattern Extraction Process:**

Episodes in Memory

```

|           |
+-----+
| 1. Cluster Similar Episodes          |
|   * Group by task type              |
|   * Group by problem domain         |
|   * Group by strategies used        |
+-----+
|           |
+-----+
| 2. Compare Within Clusters          |
|   * What's common to successes?     |
|   * What's common to failures?      |
|   * What differentiates success/fail? |
+-----+
|           |
+-----+
| 3. Identify Recurring Structures    |
|   * Frequent subsequences          |
|   * Repeated decision patterns     |
|   * Common adaptation triggers     |
+-----+
|           |
+-----+
| 4. Generalize Patterns             |
|   * Abstract from specific instances |
|   * Identify context conditions     |
|   * Estimate pattern confidence    |
+-----+
|           |
+-----+
| 5. Validate Patterns               |
+-----+

```

```

| * Statistical significance?      |
| * Consistent across episodes?   |
| * Explanable causality?         |
+-----+
|
Validated Patterns

```

## Pattern Representation:

```

Pattern {
    pattern_id: string,
    type: PatternType,
    name: string,
    description: string,

    // Context where pattern applies
    context: {
        problem_type: string,
        domain: string,
        conditions: [Condition]
    },

    // Pattern structure
    structure: {
        type: "SEQUENCE" | "CHOICE" | "CONDITION" | "OPTIMIZATION",
        elements: [...]
    },

    // Evidence
    evidence: {
        supporting_episodes: [episode_id],
        counter_episodes: [episode_id],
        confidence: float, // 0-1
        statistical_significance: float
    },

    // Application
    application: {
        when_to_apply: string,
        expected_benefit: string,
        estimated_improvement: {
            time_saving: float,
            cost_saving: float,
            success_rate_increase: float
        }
    },
}

// Metadata
discovered_at: datetime,

```

```

    last_validated: datetime,
    usage_count: int,
    success_when_applied: int
}

```

### **Example Pattern:**

Pattern: "Install-Import-Verify Sequence"

Type: SUCCESS\_PATTERN

Context: Package installation tasks in Python

Structure:

```

SEQUENCE [
  1. Try pip install
  2. If fail, check for alternatives
  3. Import package to verify
  4. Run simple test to confirm functionality
]

```

Evidence:

- \* Applied in 45 episodes
- \* Success rate: 92% (vs 70% without pattern)
- \* Statistical significance:  $p < 0.01$

Application:

When: Any Python package installation task

Benefit: Higher success rate, earlier failure detection

Improvement: ~30% time saving, ~20% higher success

Last Validated: 2024-01-15

Usage Count: 45

Success Rate: 92%

## **4.5 Strategy Evaluator**

**Purpose:** Valutare effectiveness di diverse strategie e aggiornare preferenze.

### **Strategy Comparison:**

STRATEGY EVALUATION MATRIX					
		Episodes	Success	Avg Time	Avg Cost
HTN Plan	25	88%	180s	\$0.15	

	Linear	18	78%	120s	\$0.10	
	Explore	12	83%	240s	\$0.25	
	+-----+					
	Analysis:					
	* HTN: Best success rate, moderate cost					
	* Linear: Fastest but lower success					
	* Exploratory: High cost, not justified by success					
	Recommendation:					
	* Default to HTN for code refactoring					
	* Use Linear only for simple, well-understood cases					
	* Avoid Exploratory unless high uncertainty					
	+-----+					

## Strategy Learning Process:

For each strategy type:

	+-----+					
	1. Collect Usage Data					
	* # times used					
	* Contexts used					
	* Outcomes					
	+-----+					
	+-----+					
	2. Compute Performance Metrics					
	* Success rate					
	* Average time					
	* Average cost					
	* Quality of results					
	+-----+					
	+-----+					
	3. Identify Contexts					
	* Where does it work best?					
	* Where does it fail?					
	* What conditions affect it?					
	+-----+					
	+-----+					
	4. Compare to Alternatives					
	* Which strategy is best when?					
	* Trade-offs between strategies					
	+-----+					
	+-----+					
	5. Update Strategy Preferences					

```

| * Adjust default choices      |
| * Update selection heuristics |
| * Refine context-strategy mapping |
+-----+

```

## 4.6 Performance Analyzer

**Purpose:** Analizzare performance effettiva vs aspettativa, identificare bottleneck.

### Analysis Types:

```

+-----+
|           PERFORMANCE ANALYSIS TYPES          |
|
| 1. ESTIMATION ACCURACY
|   Compare estimated vs actual:
|     * Time per task
|     * Cost per task
|     * Resource usage
|     -> Improve estimation models
|
| 2. BOTTLENECK IDENTIFICATION
|   Find tasks that:
|     * Took disproportionately long
|     * Blocked other tasks
|     * Consumed excessive resources
|     -> Target for optimization
|
| 3. EFFICIENCY OPPORTUNITIES
|   Identify:
|     * Redundant work
|     * Missed parallelization
|     * Suboptimal tool choices
|     -> Apply optimizations
|
| 4. RESOURCE UTILIZATION
|   Analyze:
|     * Token efficiency (output quality per token)
|     * Model selection appropriateness
|     * Tool usage patterns
|     -> Optimize resource allocation
+-----+

```

### Bottleneck Detection:

Analysis Process:

1. IDENTIFY SLOW TASKS
 

```
FOR task IN executed_tasks:
```

```
IF task.actual_time > task.estimated_time * 2:  
    FLAG as bottleneck candidate
```

## 2. ANALYZE ROOT CAUSE

```
FOR bottleneck IN candidates:  
    * Was it LLM inference time?  
    * Was it tool execution time?  
    * Was it waiting for dependencies?  
    * Was it due to retries/failures?
```

## 3. ASSESS IMPACT

```
* Did it block other tasks?  
* How much did it extend total time?  
* Could it be optimized?
```

## 4. GENERATE RECOMMENDATIONS

```
* Use smaller model if quality acceptable  
* Cache results if repeated  
* Parallelize if dependencies allow  
* Choose different tool/approach
```

## **Example Performance Analysis:**

Episode: "Refactor Authentication Module"

Estimated Total Time: 300s

Actual Total Time: 450s

Efficiency: 0.67 (below target of 0.8)

Breakdown:

Task	Est	Actual	Variance
Research JWT libs	30s	28s	-7% [v]
Design token flow	45s	52s	+16%
Install library	10s	35s	+250%
Create encoder	60s	65s	+8% [v]
Create decoder	60s	58s	-3% [v]
Update auth flow	90s	145s	+61%
Write tests	120s	115s	-4% [v]
Run tests	15s	12s	-20% [v]

BOTTLENECKS IDENTIFIED:

### 1. Install Library (+25s)

Root Cause: Package name issues, multiple retries

Impact: 25s delay, 17% of overrun

Recommendation: Improve package search, better contingencies

2. Update Auth Flow (+55s)
  - Root Cause: Integration complexity underestimated
  - Impact: 55s delay, 37% of overrun
  - Recommendation: Better complexity estimation for integration tasks

EFFICIENCY OPPORTUNITIES:

1. Encoder & Decoder could be parallelized
  - Current: 65s + 58s = 123s sequential
  - Potential: max(65s, 58s) = 65s parallel
  - Saving: 58s (13% of total time)
2. Library research could be cached
  - If similar task recurs, save 28s

UPDATED ESTIMATES:

Next similar task estimate adjusted:

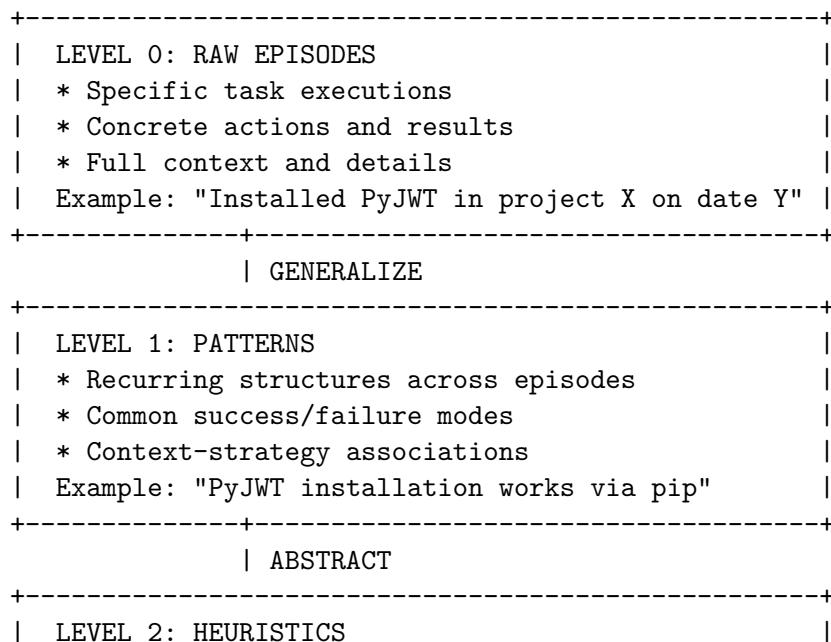
- Install library: 30s (was 10s)
  - Update auth flow: 120s (was 90s)
  - Apply parallelization pattern
- > New estimate: 280s (vs original 300s, actual 450s)

## 4.7 Knowledge Distiller

**Purpose:** Generalizzare da episodi specifici a conoscenza riusabile.

**Distillation Process:**

Specific Experiences -> General Knowledge



```

| * Rules of thumb
| * Decision guidelines
| * Optimization principles
| Example: "For JWT in Python, prefer PyJWT"
+-----+
    | FORMALIZE
+-----+
| LEVEL 3: STRATEGIES
| * Reusable approaches
| * Template solutions
| * Meta-knowledge
| Example: "Token-based auth pattern for APIs"
+-----+
    | INTERNALIZE
+-----+
| LEVEL 4: INTUITIONS
| * Implicit preferences
| * Automatic choices
| * Expertise
| Example: "Auth refactoring -> security first"
+-----+

```

## **Heuristic Generation:**

```

Heuristic {
  heuristic_id: string,
  category: "SELECTION" | "ORDERING" | "OPTIMIZATION" | "SAFETY",
  description: string,

  // When to apply
  applicability: {
    problem_types: [string],
    conditions: [Condition],
    confidence_required: float
  },

  // What to do
  guidance: {
    type: "PREFER" | "AVOID" | "ALWAYS" | "NEVER" | "IF_THEN",
    action: string,
    rationale: string
  },

  // Supporting evidence
  evidence: {
    patterns: [pattern_id],
    success_rate: float,
    episodes: [episode_id]
  }
}

```

}

### **Example Heuristics:**

Heuristic 1: "Verify Before Commit"

Category: SAFETY

Applicability: Code modification tasks

Guidance: ALWAYS run tests before committing code changes

Rationale: Prevents broken commits, catches issues early

Evidence: 95% of successful code changes included pre-commit testing

80% of failed deployments lacked pre-commit testing

Heuristic 2: "Small Model for Simple Tasks"

Category: OPTIMIZATION

Applicability: Tasks with complexity=LOW, no creativity needed

Guidance: PREFER small/fast model over large model

Rationale: Cost/latency optimization without quality loss

Evidence: 89% of simple tasks succeeded with small model

Cost saving: 90%, Time saving: 70%

Heuristic 3: "Parallel Independence"

Category: OPTIMIZATION

Applicability: Tasks with no data dependencies

Guidance: IF no dependencies THEN parallelize

Rationale: Linear speedup for independent tasks

Evidence: 2-4x speedup observed in 78% of opportunities

Heuristic 4: "Backup Before Destructive"

Category: SAFETY

Applicability: Tasks with irreversible operations

Guidance: ALWAYS create backup/checkpoint before destructive ops

Rationale: Enables rollback if things go wrong

Evidence: 100% of recovered failures had backups

0% recovery rate without backups

## **4.8 Memory Updater**

**Purpose:** Persistere insights dalla reflection in memory systems.

### **Update Operations:**

MEMORY UPDATE OPERATIONS	
1. EPISODIC MEMORY UPDATE	
* Store complete episode	
* Add metadata (tags, category, outcome)	
* Create embeddings for semantic search	
* Link to related episodes	

2. PATTERN CACHE UPDATE
* Add newly discovered patterns
* Update existing pattern statistics
* Retire low-confidence patterns
* Consolidate similar patterns
3. STRATEGY PREFERENCES UPDATE
* Adjust strategy selection weights
* Update context-strategy mappings
* Refine estimation models
4. HEURISTICS UPDATE
* Add new heuristics
* Strengthen validated heuristics
* Weaken contradicted heuristics
* Remove obsolete heuristics

### **Memory Consolidation:**

Periodic Consolidation Process:

1. MERGE SIMILAR EPISODES
  - \* Find highly similar episodes
  - \* Abstract common structure
  - \* Keep representative + summary
  - \* Save storage space
2. STRENGTHEN VALIDATED PATTERNS
  - \* Patterns confirmed by new episodes
  - \* Increase confidence score
  - \* Promote to higher priority
3. PRUNE LOW-VALUE CONTENT
  - \* Patterns with low confidence
  - \* Episodes very old and never retrieved
  - \* Heuristics never applied
  - > Archive or delete
4. RESOLVE CONFLICTS
  - \* Patterns that contradict
  - \* Heuristics that conflict
  - > Keep higher-evidence version
5. GENERALIZE WHERE POSSIBLE
  - \* Multiple specific patterns
  - > Combine into general pattern

## 4.9 Reflection Output Schema

```
ReflectionInsights {
    // Episode Summary
    episode_id: string,
    episode_summary: {
        task_type: string,
        outcome: "SUCCESS" | "PARTIAL" | "FAILURE",
        key_events: [Event],
        duration: float,
        cost: float
    },
    // Performance Analysis
    performance: {
        efficiency_ratio: float,
        bottlenecks: [
            {
                task_id: string,
                delay: float,
                root_cause: string,
                recommendation: string
            }
        ],
        estimation_errors: {
            time_error: float, // actual/estimated
            cost_error: float
        }
    },
    // Patterns Discovered
    new_patterns: [Pattern],
    pattern_validations: [
        {
            pattern_id: string,
            outcome: "VALIDATED" | "CONTRADICTED",
            evidence: string
        }
    ],
    // Strategy Insights
    strategy_evaluations: [
        {
            strategy: string,
            effectiveness: float,
            context: string,
            recommendation: "INCREASE_USE" | "DECREASE_USE" | "REFINE"
        }
    ]
}
```

```

] ,  

// Learned Heuristics  

new_heuristics: [Heuristic],  

heuristic_updates: [  

  {  

    heuristic_id: string,  

    confidence_delta: float,  

    reason: string  

  }  

],  

// Improvement Opportunities  

opportunities: [  

  {  

    type: "OPTIMIZATION" | "QUALITY" | "COST_REDUCTION" | "SPEED",  

    description: string,  

    potential_impact: string,  

    implementation_suggestion: string  

  }  

],  

// Failures & Lessons  

failures: [  

  {  

    what_failed: string,  

    why_failed: string,  

    how_to_prevent: string,  

    generalizability: float  

  }  

],  

// Questions & Uncertainties  

open_questions: [  

  {  

    question: string,  

    why_uncertain: string,  

    how_to_resolve: string  

  }  

],  

// Memory Updates  

memory_updates: {  

  episodic_memory: "STORED",  

  patterns_added: int,  

  patterns_updated: int,  

  heuristics_added: int,  

  heuristics_updated: int
}

```

```

    },
    // Meta-Learning
    meta_insights: [
      {
        insight: string,
        category: "LEARNING_STRATEGY" | "ADAPTATION" | "LIMITS",
        implication: string
      }
    ]
}

```

## 5. Cognitive Layer Integration

### 5.1 Complete Flow Example

**Scenario:** User requests “Add rate limiting to API endpoints”

+-----+	
PHASE 1: GOAL ANALYSIS	
Input: "Add rate limiting to API endpoints"	
Semantic Parsing:	
Intent: MODIFICATION (adding feature)	
Entities: ["rate limiting", "API endpoints"]	
Actions: ["add", "implement"]	
Goal Extraction:	
Primary: Implement rate limiting on API	
Sub-goals:	
1. Choose rate limiting strategy	
2. Select rate limiting library/approach	
3. Implement rate limiter	
4. Apply to API endpoints	
5. Test rate limiting behavior	
Constraints:	
* Must not break existing endpoints	
* Should be configurable	
* Must handle distributed systems (implied)	
Context:	
* Codebase: Python FastAPI application	
* Similar past task: "Add authentication" (retrieved)	
* Pattern: Middleware pattern applicable	
Complexity: MODERATE (5 subtasks, moderate uncertainty)	

```

| Strategy: HTN Planning recommended |
+-----+
|
+-----+
| PHASE 2: PLANNING |
|
| Strategy Selected: HTN Planning |
| |
| Task Decomposition: |
| L0: Add rate limiting |
| L1: |
|     1.1 Research rate limiting approaches |
|     1.2 Design rate limiting architecture |
|     1.3 Implement rate limiter |
|     1.4 Integrate with FastAPI |
|     1.5 Test and validate |
| L2: |
|     1.1.1 Search for FastAPI rate limiting libraries |
|     1.1.2 Compare options (slowapi, fastapi-limiter) |
|     1.1.3 Choose best fit |
|     1.3.1 Install chosen library |
|     1.3.2 Configure rate limits |
|     1.4.1 Add middleware to app |
|     1.4.2 Apply decorators to endpoints |
|     1.5.1 Write test cases |
|     1.5.2 Run tests |
|     1.5.3 Manual verification |
| |
| Dependency Analysis: |
|     1.1 -> 1.2 -> 1.3 -> 1.4 -> 1.5 (mostly sequential) |
|     1.5.1 -> 1.5.2 (sequential within test phase) |
| |
| Resource Estimates: |
|     Total: ~240s, $0.18, 85K tokens |
| |
| Contingencies: |
|     * If library install fails -> try alternative |
|     * If tests fail -> debug and retry |
+-----+
|
+-----+
| PHASE 3: EXECUTION |
|
| Task 1.1.1: Search for FastAPI rate limiting libraries |
|     -> Execute: Web search + document read |
|     -> Result: Found [slowapi, fastapi-limiter, ...] |
|     -> Verification: [v] PASS |
| |

```

```
| Task 1.1.2: Compare options
|   -> Execute: LLM analysis of options
|   -> Result: "slowapi recommended (simple, Redis-backed)"
|   -> Verification: [v] PASS
|
| Task 1.1.3: Choose best fit
|   -> Execute: Decision based on analysis
|   -> Result: "slowapi selected"
|   -> Verification: [v] PASS
|
| Task 1.3.1: Install chosen library
|   -> Execute: pip install slowapi
|   -> Result: FAILURE "Package not found"
|   -> Apply Contingency: Try "pip install slowapi"
|   -> Result: SUCCESS (typo correction worked)
|   -> Verification: [v] PASS
|
| Task 1.3.2: Configure rate limits
|   -> Execute: Write configuration code
|   -> Result: Configuration created
|   -> Verification: [v] PASS
|
| ... (continues through all tasks)
|
| Final Result: SUCCESS
|   * All tests passing
|   * Rate limiting working as expected
|   * Documentation updated
```

```
+-----+
| |
+-----+
| PHASE 4: REFLECTION
|
| Episode Analysis:
|   * Total time: 265s (vs estimated 240s) = 1.1x
|   * Total cost: $0.19 (vs estimated $0.18) = 1.05x
|   * Success rate: 100% (1 retry needed)
|   * Plan adherence: 95%
|
| Patterns Discovered:
|   * "slowapi" is preferred library for FastAPI
|     rate limiting (add to pattern cache)
|   * Middleware pattern worked well
|     (validate existing pattern)
|
| Performance Insights:
|   * Installation phase took longer due to typo
|   * Testing was faster than estimated
```

```

|   * Overall efficiency: 0.9 (good)
|
| Learned Heuristic:
|   "For FastAPI features, check if middleware
|     approach is applicable before custom implementation"
|
| Memory Updates:
|   * Episode stored in episodic memory
|   * Pattern "FastAPI-rate-limiting" added
|   * Estimation model for API features refined
|
| Opportunities Identified:
|   * Could cache library search results
|   * Could create template for API feature additions
+-----+

```

## 5.2 Caratteristiche del Cognitive Layer

### Proprietà Chiave:

#### MODULARITÀ

- \* Ogni componente ha responsabilità chiare
- \* Interfacce pulite tra componenti
- \* Componenti migliorabili indipendentemente

#### BOUNDED EMERGENCE

- \* Reasoning LLM entro framework strutturato
- \* Validazione esplicita ad ogni stadio
- \* Bounds di sicurezza enforced ovunque

#### APPRENDIMENTO & ADATTAMENTO

- \* Apprende da ogni esecuzione
- \* Adatta strategie basate sull'esperienza
- \* Costruisce conoscenza riusabile nel tempo

#### TRASPARENZA

- \* Trace di esecuzione complete
- \* Reasoning esplicito catturato
- \* Decisioni spiegabili

#### ROBUSTEZZA

- \* Multipli layer di validazione
- \* Degradazione graduale
- \* Recovery automatico da errori quando possibile

## 5.3 Caratteristiche di Performance

### Breakdown Latenza (task complesso tipico):

Analisi Goal: 15-30s (10-15%)

Pianificazione: 30-60s (20-30%)

Esecuzione: 120-240s (50-70%)

Reflection: 20-40s (10-15%)

---

Totale: 185-370s

Nota: Reflection è spesso asincrona,  
non blocca risposta utente

### **Utilizzo Risorse:**

Token:

- \* Analisi Goal: 5-10K
- \* Pianificazione: 10-20K
- \* Esecuzione: 40-80K (varia per task)
- \* Reflection: 8-15K

Totale: 63-125K token per task complesso

Chiamate LLM:

- \* Analisi Goal: 2-3
- \* Pianificazione: 3-5
- \* Esecuzione: 10-30 (dipende da task)
- \* Reflection: 2-4

Totale: 17-42 chiamate per task complesso

Costo:

- \* \$0.15-0.35 per task complesso (dipende da modello)

---

**Prossimo:** 03-memory-system.md -> Specifiche dettagliate dell'architettura di memoria